

Proving MCAPI Executions are Correct

Applying SMT Technology to Message Passing

Yu Huang, Eric Mercer, and Jay McCarthy *

Brigham Young University
{yuHuang, egm, jay}@byu.edu

Abstract. Asynchronous message passing for C is an important paradigm in writing applications for embedded heterogeneous multicore systems. The Multicore Association (MCA), an industry consortium promoting multicore technology, is working to standardize message passing into a single API, MCAPI, for bare metal implementation and portability across platforms. Correctness in such an API is difficult to reason about manually, and testing against reference solutions is equally difficult as reference solutions implement an unknown set of allowed behaviors, and programmers have no way to directly control API internals to expose or reproduce errors. This paper encodes an MCAPI execution as an SMT problem, which if satisfiable, yields a feasible execution schedule on the same trace, such that it resolves non-determinism in the MCAPI runtime in such a way that it now fails user provided assertions. Such an encoding is useful for test, debug, and verification of MCAPI program execution. The novelty in the encoding is the use of match-pairs (potential send and receive couplings) to model the MCAPI execution in the SMT problem. Match-pair encoding for MCAPI executions, when compared to other encoding strategies, is simpler to reason about, results in significantly fewer terms in the SMT problem, and captures feasible behaviors that are ignored in previously published techniques.

Keywords: Abstraction, refinement, SMT, message passing

1 Introduction

Embedded devices fill all sorts of crucial roles in our lives. They exist as medical devices, as network infrastructure, and they control our automobiles. Embedded devices continue to become more powerful as computing hardware becomes smaller and more modular. It is now becoming commonplace to find multiple processing units inside a single device. The Multicore Association (MCA) is an industry group that has formed to define specifications for low-level communication, resource management, and task management for multicore devices.

One specification that the MCA has released is the Multicore Association Communications API (MCAPI). The specification defines basic data type, data structures, and functions that can be used to perform simple message passing

* Special thanks to Christopher Fischer, formally at BYU and now at Amazon

operation between different computing entities within a device. Messages can be passed across persistent channels that force an ordering of the messages, or they can be passed to specific *endpoints* within the system. The specification places few ordering constraints on messages passed from one endpoint to another. This freedom introduces into the system possibilities for non-deterministic behavior in the way the messages arrive at their destination.

This non-deterministic behavior in MCAPI programs is difficult for current tools to test or debug. Chao Wang provided a description of a verification method for multithreaded software that uses shared memory as a means of synchronization that provides a marked improvement over a comparable tool that uses concrete execution [4]. Unfortunately, Wang’s method for shared memory system does not work on the message passing system. Sharma et al. created a method of using concrete execution to verify MCAPI programs, but it was later discovered that this method does not completely explore the entire execution space of certain kinds of MCAPI programs [3]. Elwakil et al. also defined a method of representing MCAPI program executions as satisfiability problems building on the method of [4] adapting it to message passing, but this method does not correctly model all kinds of MCAPI program executions [2].

Either the existing technique can not correctly encodes the program, or they can not provide efficient method for testing or debugging. Due to the difficulty and defect described above, we present an encoding for MCAPI program executions that uses match-pairs from [3] in a new SMT encoding that requires fewer terms than [2] but includes all possible program executions unlike [2]. In other words, the encoding itself provides efficiency. Most importantly, this method guarantees of finding all hidden errors by encoding all possible execution traces of a concurrent program into an SMT problem.

Our main contributions in this paper include:

1. A correct and efficient SMT encoding of an MCAPI program execution; and
2. An $O(n^2)$ algorithm to generate an over-approximation of possible match-pairs.

The rest of the paper is organized as follows: Section 2 presents an MCAPI program, where two outcomes of execution traces can be found. Section 3 defines an SMT encoding for an MCAPI program execution that captures non-determinism in message order by defining the set of possible match pairs between sends and receives on a common endpoint. Section 4 solves the outstanding problem in other encodings of generating feasible match-pairs to use in the encoding. We present a $O(n^2)$ algorithm that over-approximates the true match set. We then use our encoding in a *CEGAR*¹ loop to prune infeasible match pairs from the initial set until we witness a true error or conclude no error is possible. Section 5 presents the experimental results that show our encoding to be correct and efficient, as other existing encodings both omit critical program behaviors and require more SMT clauses. Section 6 proves that our encoding is correct by

¹ The *CEGAR* loop is defined as a **C**ounter **E**xample **G**uided **A**bstraction **R**efinement.

first giving formal semantics to an MCAPI program execution. Second, defining the encoding of a single program execution in an SMT program; and third, generalizing that encoding to set of match pairs rather than those generated in the specific program encoding. Section 7 discusses related work, and Section 8 presents our conclusions.

2 Example

It is a challenge to explain intended behavior in simple scenarios consisting of a handful of calls when dealing with concurrency. Consider the MCAPI program in the top of Figure 1 that includes three tasks that use send (`mcapi_msg_send_i`) and receive (`mcapi_msg_rcv_i`) calls to communicate with each other. Line numbers appear in the second column for each task, and the declarations of the local variables are omitted for space. Picking up the scenario just after the endpoints are defined, lines 02 and 05 of task 0 receive two messages on the endpoint *ep0* in variables *A* and *B* which are converted to integer values and stored in variables *a* and *b* on lines 04 and 07; task 1 receives one message on endpoint *ep1* in variable *C* on line 03 and then sends the message “1” on line 05 to *ep0*; and finally, task 2 sends messages “4” and “Go” on lines 04 and 06 to endpoints *ep0* and *ep1* respectively. Task 0 has additional code (lines 08 - 10) to assert properties of the values in *a* and *b*. The `mcapi_wait` calls block until the associated send or receive buffer is able to be used. Given the scenario, we might ask the question: “What are the possible values of *a* and *b* after the scenario completes?”

Intuitively, variable *a* should contain 4 and variable *b* should contain 1 since task 2 must first send message “4” to *ep0* before it can send message “Go” to *ep1*; consequently, task 1 is then able to send message “1” to *ep0*. At the end of execution the assertion on line 12 of task 0 holds and no error is found. Such intuition is a valid program execution, and the sequential trace (ignoring endpoint creation) is given in the first column for each task in Figure 1 starting on line 04 of task 2. This single sequential execution is not sufficient to prove correctness however because of concurrency as other valid traces may exist. How to test those other traces is the question answered in this paper.

A single execution through a concurrent program can be understood as a concurrent trace program (CTP) which is a trace divided into its concurrent components including the requisite control flow on which the trace is valid [4]. The CTP for the trace indicated in the first column of each task in Figure 1 is shown in the bottom of the same figure. The CTP is more abstract than the original C code, and line numbers are now grouped into a program location (where the number in front of “.” is the thread number and the number behind is the command order) in each command, but it retains sufficient details to reason about other possible executions. The assume statement on program location 0_4 of task 0 preserves the control flow of the original trace. In other words, any alternative trace on the CTP is only valid if the assume statement holds; otherwise, such an execution belongs to a different path through the program. We define a trace language that takes as input a single execution trace and the

CTP, to verify the correctness of executing a single trace. The trace language simulates the trace execution by rewriting the program state based on transitions described in the trace and the CTP. We apply queue movement in the language semantic to specify the communication between a send and a receive command. By applying the same framework of our trace language, an SMT problem can be defined on the simulation of a trace execution. An SMT solver for the defined problem is used to enumerate the possible valid executions in an effort to find one that violates the assertion on program location 0.5 of the CTP for our example while preserving the assumption on line 0.4.

The novelty of the SMT encoding in this paper is its use of match pairs rather than the state-based or order-based encoding of prior work [2, 1]. A match pair is the coupling of a receive to a particular send. Figure 2(a) is the set of possible match pairs for the CTP in the bottom of Figure 1 for our running example. The set admits, for example, that `rcvA` can be matched with either `snd1` or `snd2`. The SMT encoding in this paper asks the SMT solver to resolve the match pairs for the system.

Prior SMT models of MCAPI program executions implicitly compute match-pairs by adding possible happens-before relations on sends and receives with the conditions under which those orderings are valid. The SMT solver is then asked to resolve the happens-before relation by choosing specific orders of sends and receives. Match-pair encoding, as presented in this work, though it has the same computational complexity as order based encoding (you still need to figure out match pairs on endpoints), is simpler to reason about directly rather than implicitly through orders, results in significantly fewer terms in the SMT problem, and does not restrict out possible matches in non-blocking actions—the order based encoding misses valid traces.

Figure 2(b) is the stylized SMT encoding for the CTP from our running example. Not shown are the definitions for `rcvA`, `rcvB`, etc. that are not novel to our solution but are records that group the various parameters for the respective send and receive calls with an event term to encode order. In our simple example, the send definitions have the match-pair and event term unconstrained. Receive definitions leave the event term, match-pair, and final value of the receive variable unconstrained. All of these unconstrained variables must be resolved by the SMT solver. The variables ending in “`_loc`” represent the event variables for the corresponding MCAPI operations. The `HB` terms encode constraints on the event variables to enforce program order as it appears in each task of the CTP. For example, by program order, `rcvA` happens before the wait command `wait_rcvA` for `rcvA` as shown in Figure 2(b). Similarly, `wait_rcvA` happens before `rcvB`. Any assignment to event variables must comply with program order constraints.

The `MATCH` terms in Figure 2(b) add constraints requisite with the specified match pair: the send and receive endpoints must be the same, the send value must appear in the receive variable, the send and receive events must be ordered appropriately, and the two events must be matched together. For example, matching `rcvA` with `snd3` forces the variable `a` to have the value 1. The `MATCH` function encodes this relationship. The `ne` term prevents the SMT solver

from finding solutions that use conflicting match pairs. In our example, the NE term precludes `rcvA` and `rcvB` from matching to the same send. The first assert works as the assume in Figure 1, such that it prevents the SMT solver from finding solutions that are not consistent with control flow, so any solution that has $b \leq 0$ is infeasible. Finally, the second assert is negated as the goal is to find schedules that violate the property. In our example, the interest is in solutions that resolve in $a \neq 4$.

The SMT encoding in Figure 2(b) can be given to any preferred SMT solver and a satisfying assignment is a feasible execution that violates an assertion. For the example in Figure 1 such an execution does exist. The MCAPI natural language description states the send operation “*returns once the buffer can be reused by the application.*”. As such, the return of the send only implies a copy-out of the message buffer and not a delivery to the intended endpoint; thus, an additional program execution places the value 1 in variable a and the value 4 in variable b . In order to place the value of some send to a variable of a specific receive call, we apply the queue movement in the sequential execution, where the pending sends in the source task are moved to the send queue in the destination task. The first column of each task in the bottom of Figure 2 gives the sequential program order of each statement for the error trace, and the second column indicates the appropriate queue movements where the match pairs are implied in Figure 2. For example, step 09 of the trace (found in task 0) indicates that the first item of the pending sends in task 1 is moved to the send queue in task 0, and the wait command on `rcvA` in program location 0_1 matches with `snd3` (the first item in the send queue in task 0). Conversations with the MCAPI designers confirm the intended behavior of the API to include both program executions of the scenario. To date, there have been three published verification and analysis tools purpose-built for MCAPI that omit this less intuitive program execution [3, 2, 1].

Finally, since the computation of match-pairs is expensive to do precisely, our solution includes the ability to over-approximate the match-pair set, and then validate satisfying assignments from the SMT solver to see if they use only feasible match pairs. Such validation is accomplished by creating a model of the MCAPI library with full operational semantics. The model is able to detect bad match pairs in a given program trace. The bad match-pairs can be omitted from the SMT encoding and the SMT solver re-run to see if another solution exists. The framework provides a *CEGAR* loop for MCAPI trace verification. The remainder of this paper details our MCAPI trace model, encoding, and an algorithm to compute an over-approximated match pair set.

3 Trace Language

The trace language is the theoretical framework for the match-pair encoding. The language syntax describes a CTP with a single execution trace on the same CTP. The evaluation syntax with its operational semantics define how to execute the CTP, following the specified trace, and define when that execution is a

Task 0	Task 1	Task 2
<pre> 00 initialize(NODE_0,&v,&s); 01 ep0 = create_endpoint(PORT_0,&s); 04 02 msg_rcv_i(ep0,A,sizeof(A),&rcvA,&s); 05 03 wait(&rcvA,&size,&s,MCAPI_INF); 06 04 a = atoi(A); 11 05 msg_rcv_i(ep0,B,sizeof(B),&rcvB,&s); 12 06 wait(&rcvB,&size,&status,MCAPI_INF); 13 07 b = atoi(B); 14 08 if (b > 0) 15 09 assert(a == 4); 10 finalize(&s); </pre>	<pre> 00 initialize(NODE_1,&v,&s); 01 ep1 = create_endpoint(PORT_1,&s); 02 ep0 = get_endpoint(NODE_0,PORT_0,&s); 07 03 msg_rcv_i(ep1,C,sizeof(C),&rcvC,&s); 08 04 wait(&rcvC,&size,&s,INF); 09 05 msg_snd_i(ep1,ep0,"1",2,N,&snd3,&s); 10 06 wait(&snd3,&size,&s,MCAPI_INF); 07 finalize(&s); </pre>	<pre> 00 initialize(NODE_2,&v,&s); 01 t2 = create_endpoint(PORT_2,&s); 02 t0 = get_endpoint(NODE_0,PORT_0,&s); 03 t1 = get_endpoint(NODE_1,PORT_1,&s); 00 04 msg_snd_i(ep2,ep0,"4",2,N,&snd1,&s); 01 05 wait(&snd1,&size,&status,MCAPI_INF); 02 06 msg_snd_i(ep2,ep1,"Go",3,N,&snd2,&s); 03 07 wait(&snd2,&size,&status,MCAPI_INF); 08 09 finalize(&s); </pre>
<pre> 08 0_0 (0_0 (rcv_i rcvA 0 a)) 09 (0_1 (0 1)) (0_1 (wait rcvA)) 10 0_2 (0_2 (rcv_i rcvB 0 b)) 11 (0_3 (0 2)) (0_3 (wait rcvB)) 12 0_4 (0_4 (assume (> b 0))) 13 0_5 (0_5 (assert (= a 4))) </pre>	<pre> 04 1_0 (1_0 (rcv_i rcvC 1 c)) 05 (1_1 (1 2)) (1_1 (wait rcvC)) 06 1_2 (1_2 (snd_i snd3 1 0 0x1)) 07 1_3 (1_3 (wait snd3)) </pre>	<pre> 00 2_0 (2_0 (snd_i snd1 2 0 0x4)) 01 2_1 (2_1 (wait snd1)) 02 2_2 (2_2 (snd_i snd2 2 1 0x476f00)) 03 2_3 (2_3 (wait snd2)) </pre>

Fig. 1. Two versions and traces of the same concurrent system with the top using MCAPI and the bottom using the trace language abstraction derived from the execution order in the first column of the program in the top.

	<pre> (HB rcvA_loc wait_rcvA_loc) (HB wait_rcvA_loc rcvB_loc) (HB rcvB_loc wait_rcvB_loc) (HB wait_rcvB_loc assume_loc) (HB assume_loc assert_loc) (HB rcvC_loc wait_rcvC_loc) (HB wait_rcvC_loc snd3_loc) (HB snd3_loc wait_snd3_loc) (HB snd1_loc wait_snd1_loc) (HB wait_snd1_loc snd2_loc) (HB snd2_loc wait_snd2_loc) (or (MATCH rcvA snd3) (MATCH rcvA snd1)) (or (MATCH rcvB snd1) (MATCH rcvB snd3)) (NE rcvA rcvB) (MATCH rcvC snd2) (assert (> b 0)) (assert (not (= a 4))) </pre>
(a)	(b)

Fig. 2. A match-pair set and SMT encoding of the system in bottom half of Figure 1. (a) The match pairs based on end-points. (b) The SMT encoding where HB creates a *happens-before* constraint, MATCH creates a match-pair constraint, and NE creates a constraint to preclude the use of conflicting match pairs.

$ \begin{aligned} ctp &::= (t \dots) \\ t &::= ([\rho \ c] \dots \perp) \\ c &::= (\mathbf{assume} \ e) \\ &\quad (\mathbf{assert} \ e) \\ &\quad (\mathbf{x} := e) \\ &\quad (\mathbf{sndi} \ a \ \alpha \ \beta \ e \ \rho) \\ &\quad (\mathbf{rcvi} \ a \ \beta \ \mathbf{x} \ \rho) \\ &\quad (\mathbf{wait} \ a) \\ m &::= \perp \mid \delta \\ \delta &::= ((\beta \ \alpha) \dots \perp) \\ trace &::= (\sigma \dots) \\ \sigma &::= (\rho \ m) \\ e &::= (\mathbf{op} \ e \ e) \\ &\quad c \\ &\quad \mathbf{x} \\ &\quad v \\ v &::= \mathbf{number} \\ &\quad \mathbf{bool} \\ \mathbf{bool} &::= \mathbf{true} \\ &\quad \mathbf{false} \\ \alpha &::= \gamma \\ \beta &::= \gamma \end{aligned} $	$ \begin{aligned} mstate &::= (h \ \eta \ A \ Pnd.s \ Pnd.r \ Q.s \ ctp \ trace \ s) \\ qstate &::= (Pnd.s \ Q.s \ m \ s) \\ estate &::= (h \ \eta \ A \ Pnd.s \ Pnd.r \ Q.s \ e \ s \ k) \\ h &::= \emptyset \mid (h \ [l \rightarrow v]) \\ \eta &::= \emptyset \mid (\eta \ [\mathbf{x} \rightarrow l]) \\ A &::= ((a \ \gamma) \dots) \\ Pnd.s &::= \emptyset \mid (Pnd.s \ [\beta \rightarrow frm]) \\ frm &::= \emptyset \mid (frm \ [\alpha \rightarrow snd]) \\ snd &::= ([a \ \mathbf{v}] \dots) \\ Pnd.r &::= \emptyset \mid (Pnd.r \ [\beta \rightarrow rcv]) \\ rcv &::= ([a \ \mathbf{x}] \dots) \\ Q.s &::= \emptyset \mid (Q.s \ [\beta \rightarrow q]) \\ q &::= ([a \ v \perp] \dots) \\ v \perp &::= \perp \mid \mathbf{v} \\ s &::= \mathbf{success} \\ &\quad \mathbf{failure} \\ &\quad \mathbf{infeasible} \\ &\quad \mathbf{error} \\ &\quad \mathbf{ret} \\ k &::= \mathbf{ret} \\ &\quad (\mathbf{assert} \ * \rightarrow k) \\ &\quad (\mathbf{assume} \ * \rightarrow k) \\ &\quad (\mathbf{x} := \ * \rightarrow k) \\ &\quad (\mathbf{op} \ * \ e \rightarrow k) \\ &\quad (\mathbf{op} \ \mathbf{v} \ * \rightarrow k) \end{aligned} $
---	--

(a)

(b)

Fig. 3. The trace language syntax with its evaluation syntax for the operational semantics—bold face indicates a terminal. (a) The input syntax with terminals \mathbf{x} , ρ (which is unique), and a defined as strings and γ as a number. (b) The evaluation syntax with terminal l defined as a number.

$$\begin{array}{l}
\text{MACHINE STEP} \\
(Pnd.s \ Q.s \ m \ s) \rightarrow_q^* (Pnd.s_{p0} \ Q.s_{p0} \ m_p \ s_{p0}) \\
(h \ \eta \ A \ Pnd.s_{p0} \ Pnd.r \ Q.s_{p0} \ c_0 \ s_{p0} \ \mathbf{ret}) \rightarrow_e^* (h_p \ \eta_p \ A_p \ Pnd.s_{p1} \ Pnd.r_p \ Q.s_{p1} \ e \ s_{p1} \ \mathbf{ret}) \\
\hline
(h \ \eta \ A \ Pnd.s \ Pnd.r \ Q.s \ (t_0 \dots ([\rho_0 \ c_0] [\rho_1 \ c_1] [\rho_2 \ c_2] \dots) t_2 \dots) ([\rho_0 \ m] \sigma_1 \dots) s) \rightarrow_m \\
(h_p \ \eta_p \ A_p \ Pnd.s_{p1} \ Pnd.r_p \ Q.s_{p1} \ (t_0 \dots ([\rho_1 \ c_1] [\rho_2 \ c_2] \dots) t_2 \dots) (\sigma_1 \dots) s_{p1})
\end{array}$$

Fig. 4. Machine Reductions (\rightarrow_m)

$$\begin{array}{l}
\text{PROCESS QUEUE MOVEMENT} \\
([a_s \ v] [a_1 \ v_1] \dots) = Pnd.s(\beta)(\alpha) \\
Pnd.s_p = [Pnd.s \mid \beta \mapsto [Pnd.s(\beta) \mid \alpha \mapsto ([a_1 \ v_1] \dots)]] \quad ([a_1 \ v_1] \dots) = Q.s(\beta) \\
Q.s_p = [Q.s \mid \beta \mapsto ([a_s \ v] [a_1 \ v_1] \dots)] \quad s_p = \begin{cases} s & \text{if } |Pnd.s(\beta)(\alpha)| > 0 \\ error & \text{otherwise} \end{cases} \\
\hline
(Pnd.s \ Q.s \ ((\beta \ \alpha) (\beta_0 \ \alpha_0) \dots \perp) s) \rightarrow_q (Pnd.s_p \ Q.s_p \ ((\beta_0 \ \alpha_0) \dots \perp) s_p)
\end{array}$$

Fig. 5. Queue Reductions (\rightarrow_q)

VARIABLE LOOKUP $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ x \ s \ k) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ h(\eta(x)) \ s \ k)$	LEFT OPERAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (op \ e_0 \ e) \ s \ k) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e_0 \ s \ (op \ * \ e \rightarrow k))$
RIGHT OPERAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ s \ (op \ * \ e \rightarrow k)) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e \ s \ (op \ v \ * \ \rightarrow k))$	BINARY OPERATION $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v_r \ s \ (op \ v_l \ * \ \rightarrow k)) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ op(v_l, v_r) \ s \ k)$
ASSUME EXPRESSIONS EVALUATION $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (assume \ e) \ s \ k) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e \ s \ (assume \ * \ \rightarrow k))$	ASSUME COMMAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ s \ (assume \ * \ \rightarrow k)) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ assume(v, s) \ k)$
ASSERT EXPRESSIONS EVALUATION $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (assert \ e) \ s \ k) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e \ s \ (assert \ * \ \rightarrow k))$	ASSERT COMMAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ s \ (assert \ * \ \rightarrow k)) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ assert(v, s) \ k)$
ASSIGN EXPRESSIONS EVALUATION $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (x := e) \ s \ k) \rightarrow_e$ $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e \ s \ (x := * \ \rightarrow k))$	ASSIGN COMMAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ s \ (x := * \ \rightarrow k)) \rightarrow_e$ $([h \mid \eta(x) \mapsto v] \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ v \ s \ k)$
SNDI COMMAND $\frac{([a_1 \ v_1] \ \dots) = Pnd_s(\beta)(\alpha) \quad Pnd_s_p = [Pnd_s \mid \beta \mapsto [Pnd_s(\beta) \mid \alpha \mapsto ([a_0 \ h(\eta(x))] \ [a_1 \ v_1] \ \dots)]]}{(h \ \eta \ ([a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{sndi} \ a_0 \ \alpha \ \beta \ x) \ s \ k) \rightarrow_e (h \ \eta \ ([a_0 \ \alpha] \ [a \ \gamma] \ \dots) \ Pnd_s_p \ Pnd_r \ Q_s \ \mathbf{true} \ s \ k)}$	
RCVI COMMAND $\frac{([a_1 \ x_1] \ \dots) = Pnd_r(\beta) \quad Pnd_r_p = [Pnd_r \mid \beta \mapsto ([a_0 \ x_0] \ [a_1 \ x_1] \ \dots)]}{(h \ \eta \ ([a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{rcvi} \ a_0 \ \beta \ x_0) \ s \ k) \rightarrow_e (h \ \eta \ ([a_0 \ \beta] \ [a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r_p \ Q_s \ \mathbf{true} \ s \ k)}$	
WAIT (SNDI) COMMAND $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{wait} \ a_s) \ s \ k) \rightarrow_e (h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ \mathbf{true} \ s \ k)$	
WAIT (RCVI) COMMAND $\frac{([a_0 \ \gamma_0] \ \dots [a_s \ \alpha] \ \dots [a_r \ \beta] \ \dots [a_1 \ \gamma_1]) = A \quad (h_p \ a_s \ Pnd_r_p \ Q_s_p \ s_p) = \text{getMarkRemove}(h, \eta, Pnd_r, Q_s, \beta, a_r, s) \quad A_p = ([a_0 \ \gamma_0] \ \dots [a_1 \ \gamma_1])}{(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{wait} \ a_r) \ s \ k) \rightarrow_e (h_p \ \eta \ A_p \ Pnd_s \ Pnd_r_p \ Q_s_p \ \mathbf{true} \ s_p \ k)}$	

Fig. 6. Expression Reductions (\rightarrow_e)

success (causes no assertion violation), a failure (causes an assertion violation), infeasible (causes an assume to not hold), or an error (uses a bogus match pair). Section 4 defines the encoding of a trace language program as an SMT problem and extends that encoding to capture a set of possible traces using match-pairs.

3.1 Syntax

Figure 3(a) is the syntax for a trace language program. This presentation uses ellipses (...) to represent zero or more repetitions, bold-face to indicate terminals, and omits commas in tuples for cleanliness. A trace language program is a CTP with a trace defining a sequential run of the CTP. The language defines a CTP (*ctp*) as a list of threads. A thread (*t*) is a list of pairs with each pair being a program location and a command. For simplicity, commands (*c*) are restricted to assume, assert, assignment, non-blocking send (*snd*) and receive (*rcv*), and wait. We model a trace as an order of executed locations identified by (ρ) and a series of queue movements that move a message from a source queue to a destination queue. The queue movement (*m*) is either a special symbol indicating no movement or a list of move commands. The move list (δ) consists of several pairs of end-points where the first and second end-points are the destination and source end-points respectively. Each pair in the move list refers to a queue movement such that the first pending send of the source end-point should be moved to the send queue of the destination end-point. The non-terminal *a* in the grammar is a unique string identifier associated with a send or receive command referred to as an action ID in the text. The wait command takes a single action ID belonging to the associated send or receive action. The non-terminals α and β are source and destination end-points respectively. The non-terminal γ is a number. The terminal **x** is any string not mentioned in the grammar definition and represents a program variable. Expressions (*e*) are defined using prefix notation over binary operators. The bottom of Figure 1 is an example *ctp* in the trace language (omitting the first and second columns and trivially grouping each thread and its commands into an appropriate list using parenthesis).

A sequential trace of a CTP in the grammar is a list of trace entries (σ). A trace entry is a pair consisting of a program location (ρ) and either a queue movement list (*m*) or a symbol (\perp). An example of a trace can be seen in the bottom of Figure 1 in the second column by following the sequential order in the first column which starts on program location 2_0 of task 2. Notice that whenever the trace reaches a wait command on a receive action, the trace includes a queue movement list where the destination is listed followed by the source. Messages are delivered from the source queue to the destination queue. In other words, the trace resolves any non-determinism in scheduling or message buffering that is present in the CTP.

3.2 Operational Semantics

The operational semantics for the trace language are given by a term rewriting system using a *CESK* style machine² only the machine is augmented to include additional structure for modeling message passing. Figure 3(b) defines the machine state and other syntax relating to evaluation.

A machine step (\rightarrow_m) in Figure 4 moves a thread forward by a single command. The rules operate on a machine state tuple ($mstate$) defined in Figure 3(b). The tuple can be partitioned into members relating to the *CESK* machine, members relating to the message passing model, and the trace status. The *CESK* machine members are *ctp* (the list of thread command sequences), η (an environment mapping a variable x to a location l), h (a store mapping a location l to a value v), and k (a continuation). Among the members of the message passing model, A is a dictionary mapping an action ID a to an end-point γ . Pnd_s is a set of send queues where each queue is uniquely identified first by the source end-point and then by the destination end-point. The queue itself holds pairs consisting of an action ID and value (a, v) . Pnd_r is a set of receive queues where each queue is uniquely identified by the destination end-point. The queue itself holds pairs consisting of an action ID and a variable (a, x) . Q_s is also a set of queues where each queue is uniquely identified by the destination end-point. The queue itself holds pairs consisting of an action ID a and either a value v or a symbol \perp . Intuitively, Pnd_s , Q_s and Pnd_r are end-point queues tracking actions with associated values (sends) or variables (receives). Both Pnd_s and Q_s store the sends. Q_s holds delivered messages that are ready to be received. A message moves from Pnd_s to Q_s through queue movement lists in the trace on the CTP.

The trace status in the $mstate$ nine-tuple is given by s which ranges over a lattice:

$$\text{success} \prec \text{failure} \prec \text{infeasible} \prec \text{error}$$

The trace status only moves monotonically up the lattice starting from success. A success trace completes the entire trace, meets all the assume statements, and does not fail an assertion. A failure trace completes the entire trace, meets all the assume statements, but fails an assertion. An infeasible trace completes the entire trace but does not meet all the assume statements. An error is a trace that does not complete.

We have several *CESK* machines that handle different aspects of the CTP and trace. The machine step moves one step on the trace. In each machine step, the queue machine processes any queue movements in that step, and the expression machine handles any expressions in the command associated with the trace step.

The *Machine Step* inference in Figure 4 matches any $mstate$ that has a thread whose first list entry matches the program location in the head of the trace. A match on the inference rewrites the $mstate$ with new entries for each member of the nine-tuple by first applying the queue reduction relation until no more

² The *CESK* machine state is represented with a **C**ontrol string, **E**nvironment, **S**ore, and **K**ontinuation.

reductions apply (all queue movements are processed in the trace entry first) and then applying the expression reduction relation until no more reductions apply (as indicated by the asterisk). Note that queue reductions perform the queue movement such that all send actions are moved to the destination send queues. After completing all queue movements, messages can be delivered in the process of expression reduction.

The queue reduction for each command of the queue movement list in the trace entry is given in Figure 5. The definition of the $qstate$ four-tuple is presented in the evaluation syntax in Figure 3(b). The symbol \perp in the queue movement m indicates that no more queue movement follows. A message moves from pending to delivered through queue movement lists. Each reduction step processes the first pair of the queue movement list and reduction steps follow until \perp is shown.

Expression reductions for each command in the trace language are given in Figure 6 and are defined over the $estate$ nine-tuple in the evaluation syntax of Figure 3(b) which includes a continuation k . The **ret** continuation indicates that nothing follows, and an asterisk in a continuation is a place holder indication where evaluation is taking place. For example, the *Assume Expressions Evaluation* creates a continuation indicating that it is first evaluating the expression in the assume command. Once that expression reduces to a value, then the *Assume Command* inference matches to check the assumptions.

The expression reductions use several helper functions. The function $op(v_l, v_r)$ applies the “op” to the left and right operands. The function `getMarkRemove` is explained in the later paragraph. The other helper functions are defined below:

$$\begin{aligned} \text{assert}(v, s) &= \begin{cases} \text{failure} & \text{if } s \prec \text{failure} \wedge \\ & v = \text{false} \\ s & \text{otherwise} \end{cases} \\ \text{assume}(v, s) &= \begin{cases} \text{infeasible} & \text{if } s \prec \text{infeasible} \wedge \\ & v = \text{false} \\ s & \text{otherwise} \end{cases} \end{aligned}$$

Note that the status only moves monotonically up the lattice as mentioned previously. The notations $h(\eta(x))$ (Sndi Command), $Pnd_r(\beta)$ (Rcvi Command), and $Pnd_s(\beta)(\alpha)$ (Sndi Command) in Figure 6 are used for lookup. For example, $Pnd_s(\beta)(\alpha)$ returns a list of pairs of action IDs and values as defined in Figure 3(b).

The *Sndi Command* and *Rcvi Command* in Figure 6 update Pnd_r and Pnd_s respectively with information to complete a message send or receive at a wait command. Consider a portion of the *Rcvi Command*:

$$\begin{aligned} Pnd_r_p &= [Pnd_r \mid \beta \mapsto ([a_0 \ x_0] \ [a_1 \ x_1] \ \dots)] \\ ([a_1 \ x_1] \ \dots) &= Pnd_r(\beta) \end{aligned}$$

Pnd_r_p is a new set, just like the old set Pnd_r , only the new set maps the destination end-point β to its contents in the old set plus the added entry to the front of the list of the action ID and variable for the receive command being

evaluated. Considering the entire rule in the figure, it also updates the action ID map, $A = ([a_0 \ \beta] \ [a \ \gamma] \ \dots)$, to include the coupling between the action ID and its destination end-point.

The function of the *Wait(SNDI) Command* rule in Figure 6 is to consume the wait command only. The function of the other rule for the wait command is more involved. The *Wait(RCVI) Command* rule is complicated because a CTP can wait on receive actions out of program order. Consider the following trace

$$(\dots (rcvA \ \perp) (rcvB \ \perp) (wait(rcvB) \ \perp) (wait(rcvA) \ \perp) \dots)$$

omitting the regular definition for the commands and the trace. It is perfectly valid to now call $wait(rcvB)$ even though it appeared after the $rcvA$ action. The *Wait(RCVI) Command* handles this very situation in function `getMarkRemove`. Consider the same trace above, we have the following structures of the *pending receive queue* and the *delivered queue* respectively,

$$Pnd_r: (\dots (\beta \rightarrow ([rcvA \ x_0] [rcvB \ x_1] \dots)) \dots)$$

$$Q_s: (\dots (\beta \rightarrow ([snd1 \ v_0] [snd2 \ v_1] \dots)) \dots)$$

Function `getMarkRemove` in the command $wait(rcvB)$ works as the following steps: first finding the pair $[rcvB \ x_1]$ in the pending receive queue; then getting the pair $[snd2 \ v_1]$ in the delivered queue; then marking the pair $[snd1 \ v_0]$ as first assigning v_0 to variable x_0 in the heap then updating the pair $[snd1 \ v_0]$ to $[snd1 \ \perp]$; finally removing the pair $[rcvB \ x_1]$ and $[snd2 \ v_1]$ in Pnd_r and Q_s , respectively.

The syntax with the operational semantics, as presented, are implemented directly in PLT Redex. PLT Redex is a language for testing and debugging semantics using term rewriting and is part of the Racket runtime.

The following definition is important to the proof of the match pair encoding in Section 4.

Definition 1. *A machine state $(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ ctp \ trace \ s \ k)$ is well formed if and only if it reduces to a final state where the CTP and trace run to completion, having matched every send and receive call, and the status is either success, failure, or infeasible:*

$$\begin{aligned} (h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ ctp \ trace \ s \ k) &\rightarrow_m^* \\ (h_p \ \eta_p \ () \ Pnd_s_p \ Pnd_r_p \ Q_s_p \ (()) \dots \ ()) &s_p \ ret \end{aligned}$$

such that

$$\begin{aligned} \forall \beta, \forall \alpha, Pnd_s_p(\beta)(\alpha) &= () \wedge \forall \beta, Pnd_r_p(\beta) = () \wedge \\ \forall \beta, Q_s_p(\beta) &= () \wedge s_p \prec \mathbf{error} \end{aligned}$$

For convenience, we define the function $status(m)$ to return the final status after reduction of a well formed machine state m .

Intuitively, a well-formed machine state completes all the transitions of send and receive calls, meaning that there are no elements in A , Pnd_s , Pnd_r and Q_s , the commands in the CTP and the execution trace are correctly executed, and the status of state never enters **error**.

$$\begin{array}{c}
\text{MACHINE STEP} \\
(Pnd_s \ Q_s \ m \ s) \rightarrow_q^* (Pnd_s_{p0} \ Q_s_{p0} \ m_p \ s_{p0}) \\
(h \ \eta \ A \ Pnd_s_{p0} \ Pnd_r \ Q_s_{p0} \ c_0 \ s_{p0} \ \mathbf{ret} \ l \ smt) \rightarrow_e^* (h_p \ \eta_p \ A_p \ Pnd_s_{p1} \ Pnd_r_p \ Q_s_{p1} \ e \ s_{p1} \ \mathbf{ret} \ l_p \ smt_{p0}) \\
smt_{p1} = \text{addHB}(smt_{p0} \ \rho_0 \ \rho_1) \\
\hline
(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (t_0 \ \dots \ ([\rho_0 \ c_0] \ [\rho_1 \ c_1] \ [\rho_2 \ c_2] \ \dots) \ t_2 \ \dots) \ ([\rho_0 \ m] \ \sigma_1 \ \dots) \ s \ l \ smt) \rightarrow_m \\
(h_p \ \eta_p \ A_p \ Pnd_s_{p1} \ Pnd_r_p \ Q_s_{p1} \ (t_0 \ \dots \ ([\rho_1 \ c_1] \ [\rho_2 \ c_2] \ \dots) \ t_2 \ \dots) \ (\sigma_1 \ \dots) \ s_{p1} \ l_p \ smt_{p1})
\end{array}$$

Fig. 7. Machine Reductions to build the SMT model of a trace language program
 $(\rightarrow_{\text{m-smt}})$

$$\begin{array}{c}
\text{ASSUME EXPRESSIONS EVALUATION} \\
(h \ \eta \ A \ Pnd_s \ Pnd_r \ (\text{assume } e) \ s \ k \ l \ (\text{defs } (\text{any } \dots))) \rightarrow_e (h \ \eta \ A \ Pnd_s \ Pnd_r \ e \ s \ (\text{assume } * \rightarrow k) \ l \ (\text{defs } (e \ \text{any } \dots))) \\
\\
\text{ASSERT EXPRESSIONS EVALUATION} \\
(h \ \eta \ A \ Pnd_s \ Pnd_r \ (\text{assert } e) \ s \ k \ l \ (\text{defs } (\text{any } \dots))) \rightarrow_e (h \ \eta \ A \ Pnd_s \ Pnd_r \ e \ s \ (\text{assert } * \rightarrow k) \ l \ (\text{defs } ((\text{not } e) \ \text{any } \dots))) \\
\\
\text{ASSIGN EXPRESSIONS EVALUATION} \\
(h \ \eta \ A \ Pnd_s \ Pnd_r \ (x := e) \ s \ k \ l \ (\text{defs } (\text{any } \dots))) \rightarrow_e (h \ \eta \ A \ Pnd_s \ Pnd_r \ e \ s \ (x := * \rightarrow k) \ l \ (\text{defs } ((= x e) \ \text{any } \dots))) \\
\\
\text{SNDI COMMAND} \\
([a_1 \ v_1] \ \dots) = Pnd_s(\beta)(\alpha) \\
Pnd_s_p = [Pnd_s \mid \beta \mapsto [Pnd_s(\beta) \mid \alpha \mapsto ([a_0 \ h(\eta(x))] \ [a_1 \ v_1] \ \dots)]] \\
\text{any}_0 = (\text{define } a :: \text{send}) \\
\text{any}_1 = (\text{and } (= (\text{select } a \ ep) \ \beta) \ (= (\text{select } a \ value) \ x)) \\
\hline
(h \ \eta \ ([a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{sndi} \ a_0 \ \alpha \ \beta \ x) \ s \ k \ l \ ((\text{any}_d \ \dots) \ (\text{any}_a \ \dots))) \rightarrow_e \\
(h \ \eta \ ([a_0 \ \alpha] \ [a \ \gamma] \ \dots) \ Pnd_s_p \ Pnd_r \ Q_s \ \mathbf{true} \ s \ k \ l \ ((\text{any}_0 \ \text{any}_d \ \dots) \ (\text{any}_1 \ \text{any}_a \ \dots))) \\
\\
\text{RCVI COMMAND} \\
([a_1 \ x_1] \ \dots) = Pnd_r(\beta) \\
Pnd_r_p = [Pnd_r \mid \beta \mapsto ([a_0 \ x_0] \ [a_1 \ x_1] \ \dots)] \quad \text{any}_0 = (\text{define } a :: \text{recv}) \\
\text{any}_1 = (\text{and } (= (\text{select } a \ ep) \ \beta) \ (= (\text{select } a \ var) \ x_0)) \\
\hline
(h \ \eta \ ([a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{rcvi} \ a_0 \ \beta \ x_0) \ s \ k \ l \ ((\text{any}_d \ \dots) \ (\text{any}_a \ \dots))) \rightarrow_e \\
(h \ \eta \ ([a_0 \ \beta] \ [a \ \gamma] \ \dots) \ Pnd_s \ Pnd_r_p \ Q_s \ \mathbf{true} \ s \ k \ l \ ((\text{any}_0 \ \text{any}_d \ \dots) \ (\text{any}_1 \ \text{any}_a \ \dots))) \\
\\
\text{WAIT (RCVI) COMMAND} \\
([a_0 \ \gamma_0] \ \dots \ [a_s \ \alpha] \ \dots \ [a_r \ \beta] \ \dots \ [a_1 \ \gamma_1]) = A \\
(h_p \ a_s \ Pnd_r_p \ Q_s_p \ s_p) = \text{getMarkRemove}(h, \eta, Pnd_r, Q_s, \beta, a_r, s) \\
A_p = ([a_0 \ \gamma_0] \ \dots \ [a_1 \ \gamma_1]) \quad (l_p \ \text{last_}a_s) = \text{getlastsend/replace}(l, a_s, \alpha, \beta) \\
\hline
(h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ (\mathbf{wait} \ a_r) \ s \ k \ l \ (\text{defs } (\text{any}_a \ \dots))) \rightarrow_e \\
(h_p \ \eta \ A_p \ Pnd_s \ Pnd_r_p \ Q_s_p \ \text{true} \ s_p \ k \ l_p \ (\text{defs } ((\text{HB } (\text{select } \text{last_}a_s) \ MP) \ (\text{select } a_s \ MP)) \ (\text{MATCH } a_r \ a_s) \ \text{any}_a \ \dots)))
\end{array}$$

Fig. 8. Expression Reductions to build the SMT model of a trace language program
 $(\rightarrow_{\text{e-smt}})$

4 SMT Model

The next step, after formally defining the operational semantics of our trace language, is to formally define a translation from the trace language into an SMT problem that can correctly models the original execution trace.

An SMT problem that encodes a trace through a CTP needs various data structures and support functions. We have intuitively introduced most of them in the example of Section 2. Before presenting the translation framework from the trace language to the SMT problem, we need to explain the data structures and support functions in detail. We give the definitions in the Yices input language.

$$\begin{aligned} &(\text{define HB} :: (- > a::int b::int \\ &\quad (\text{subtype } (c::bool) (\text{ite } (< a b) (= c \text{ true}) (= c \text{ false})))) \end{aligned} \quad (1)$$

The HB function in equation (1) represents a happens-before relationship between two events. It takes two events a and b as parameters, such that $a, b \in \mathbb{N}$. In the SMT encoding, every location is assigned an event, and HB orders those locations (i.e., program order, etc.). This function creates a constraint such that the first event must be less than the second event, indicating that the first event occurs before the second event. This function is not only used to assert the program order of statements within the same thread, but also to assert, for any matched pair, that the send operation occurs before the receive.

$$\begin{aligned} &(\text{define-type } R \text{ (record } M::int \text{ ep}::int \text{ var}::int \text{ event}::int)) \\ &(\text{define-type } S \text{ (record } MP::int \text{ ID}::int \text{ ep}::int \text{ value}::int \text{ event}::int)) \end{aligned} \quad (2)$$

We define two records in equation (2) that are essential to encoding the SMT problem. The first record R defined above represents a receive action such that $M, ep, var, event \in \mathbb{N}$. M represents a free variable that is assigned the $event$ field of the matched record S . The endpoint ep is the end point of the receive operation, var represents a variable which is assigned a value received by the receive operation, and $event$ is the event term which indicates the program order for the receive action associated with the record. The second record S defined above is a send action such that $MP, ID, ep, value, event \in \mathbb{N}$. The fields ep and $event$, as those in R , are the end point of the send operation and event term, respectively. The field $value$ is the value being sent. MP is another free variable used to assign match pairs, and it is assigned an end point of a receive operation if the receive operation is to be matched to the send operation. ID is a unique number that identifies each send operation. Note that the integer value assigned to $event$ indicates the program order. For example, action a happens before b if and only if $event_a$ for a is greater than $event_b$ for b . Also, we keep record of two match variables, M and MP , for different purposes. M is used for preventing two receive operations from being matched with the same send operation, which is defined in function NE. MP is used for ordering two receive operations with respect to the program order of two matched send operations.

$$\begin{aligned}
&(\text{define MATCH} :: (- > r::R \ s::S \ (\text{subtype } (c::\text{bool}) \ (\text{ite} \\
&\quad (\text{and } (= (\text{select } r \ \text{ep}) (\text{select } s \ \text{ep})) \\
&\quad (= (\text{select } r \ \text{var}) (\text{select } s \ \text{value})) \\
&\quad (\text{HB } (\text{select } s \ \text{event}) (\text{select } r \ \text{event})) \\
&\quad (= (\text{select } r \ M) (\text{select } s \ ID)) \\
&\quad (= (\text{select } s \ MP) (\text{select } r \ \text{event})))) \\
&\quad (= c \ \text{true}) \\
&\quad (= c \ \text{false}))))))
\end{aligned} \tag{3}$$

The MATCH function in equation (3) takes tuple R and S as parameters. It creates a number of constraints that represent the pairing of the specified send and receive operations. This function asserts that the destination end-point of the send operation is the same as the end-point used by the receive operation, that the value sent by the send operation is the same value received by the receive operation, that the send operation occurs before the receive operation, that the “match” value in the receive tuple is equal to the “ID” value in the send tuple, and that the “MP” value in the send tuple is equal to the “event” value in the receive tuple.

$$\begin{aligned}
&(\text{define NE} :: (- > r1::R \ r2::R \ (\text{subtype } (res::\text{bool}) \ (\text{ite} \\
&\quad (= (\text{select } r1 \ M) (\text{select } r2 \ M)) \\
&\quad (= res \ \text{false}) \\
&\quad (= res \ \text{true}))))))
\end{aligned} \tag{4}$$

The NE function in equation (4) is used to assert that no two receive operations are matched with the same send operation. The parameters for this function are two receive tuples. An assertion is made that the values of their “match” fields, M , are not equal. This shows that they are paired with two different send operations.

Given the definition of the data structures and the support functions, we give the following modified framework that produces a correlating SMT problem for a given execution trace. The evaluation syntax for the machine reductions to build the SMT model of a trace language program are largely those of the regular machine reductions with a few changes below

$$\begin{aligned}
mstate &::= (h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ ctp \ trace \ s \ k \ l \ smt) \\
estate &::= (h \ \eta \ A \ Pnd_s \ Pnd_r \ Q_s \ e \ s \ k \ l \ smt) \\
smt &::= (defs \ constraints \ ML) \\
defs &::= (any \ \dots) \\
constraints &::= (any \ \dots) \\
ML &::= (ML = \emptyset \mid (ML \ [\mu \rightarrow ((\mu \ \nu) \dots)]))
\end{aligned}$$

The new syntax adds the smt member to the $mstate$ and $estate$ where the “any” term in $defs$ and $constraints$ matches any structure. The lists will be filled with definitions, HB entries, MATCH entries, NE entries, etc. as defined by the SMT machine reductions. A match-list ML is a set of match pair lists uniquely identified by a receive action ID, and each list consists of a set of match pairs for

this receive action ID and a send action ID. Note that we alpha renamed so that μ and ν are all unique labels. For convenience, we define function $dom(ML)$ and $range(ML)$ that return the set of receive action IDs in the set of match pairs of ML and the set itself, respectively. The symbol l is another member added to the $mstate$ and $estate$. It records the last send operation that happens before the current one in the same thread. The $qstate$ does not contain any SMT encodings in the reductions and it is not changed.

The changes for the reduction rules are presented in Figure 7 and Figure 8. The support function $addHB(((any_d...) (any_c...)) \rho_0 \rho_1)$ adds program location to the definition list and adds a happens before relation

$$\equiv (((define \rho_0 :: int) any_d...)((HB \rho_0 \rho_1) any_c...))$$

The reduction rules in Figure 7 and Figure 8 add constraints to the SMT problem with respect to the semantic definition of the trace language. The Machine Step in Figure 7 adds the program order of actions in the same thread to the SMT problem. The evaluation of assume, assert and assign in Figure 8 adds assertions to the SMT problem. The $Sndi$ and $Rcvi$ commands define the send and receive actions in the SMT problem, and the $Wait(Rcvi)$ reduction adds the "match" constraints to the SMT problem. Note that the function $getlastsend/replace$ picks up the last send operation, $last_{as}$, that happens before the current one in the same thread, and adds a HB relation between the MP of $last_{as}$ and that of the current send operation. The added relation ensures the messages from a common endpoint are non-overtaking (i.e., FIFO ordered).

The following definition supports the rest of the section.

Definition 2. A machine state $(h \eta A Pnd.s Pnd.r ctp trace s k)$ is *smt-enabled* if it is well formed. The SMT problem is taken from the final SMT reduced state:

$$(h \eta A Pnd.s Pnd.r ctp trace s k ((())) \rightarrow_{m-smt}^* (h_p \eta_p A_p Pnd.s_p Pnd.r_p ctp_p trace_p s_p k_p smt))$$

For convenience, we define the function $getSMT(m) \mapsto smt$ to return the SMT problem in the final state, and $ANS(getSMT(m)) \mapsto \{SAT, UNSAT\}$ to return the status of the SMT problem in the final state of m . Also, we define the relation of the range of function ANS such that $UNSAT > SAT$.

Figure 9 shows the SMT problem (minus the event definitions) generated from our new *CESK SMT* encoding machines for the program in Figure 1 on two different traces (parts (a) and (b)). Note that the encodings in Figure 9 (a) and (b) are identical except for the MATCH clauses, which are generated from the trace parts of the program. Furthermore, the encoding in Figure 9(a) violates the properties of interest encoded in the assert clauses, and the answer for the encoding in Figure 9 (a) is SAT. From this observation, we present the following lemma that indicates the equivalence between the status of our *CESK* machines and the satisfiability of the generated SMT encoding.

(HB rcvA (wait rcvA))	(HB rcvA (wait rcvA))	(HB rcvA (wait rcvA))
(HB (wait rcvA) rcvB)	(HB (wait rcvA) rcvB)	(HB (wait rcvA) rcvB)
(HB rcvB (wait rcvB))	(HB rcvB (wait rcvB))	(HB rcvB (wait rcvB))
(HB (wait rcvB) assume)	(HB (wait rcvB) assume)	(HB (wait rcvB) assume)
(HB assume assert)	(HB assume assert)	(HB assume assert)
(HB rcvC (wait rcvC))	(HB rcvC (wait rcvC))	(HB rcvC (wait rcvC))
(HB (wait rcvC) snd3)	(HB (wait rcvC) snd3)	(HB (wait rcvC) snd3)
(HB snd3 (wait snd3))	(HB snd3 (wait snd3))	(HB snd3 (wait snd3))
(HB snd1 (wait snd1))	(HB snd1 (wait snd1))	(HB snd1 (wait snd1))
(HB (wait snd1) snd2)	(HB (wait snd1) snd2)	(HB (wait snd1) snd2)
(HB snd2 (wait snd2))	(HB snd2 (wait snd2))	(HB snd2 (wait snd2))
(MATCH rcvA snd1)	(MATCH rcvA snd3)	(or (MATCH rcvA snd1)
(MATCH rcvB snd3)	(MATCH rcvB snd1)	(MATCH rcvA snd3))
(MATCH rcvC snd2)	(MATCH rcvC snd2)	(or (MATCH rcvB snd1)
		(MATCH rcvC snd3))
		(MATCH rcvC snd2)
		(NE rcvA rcvB)
(assert (> b 0))	(assert (> b 0))	(assert (> b 0))
(assert (not (= a 4)))	(assert (not (= a 4)))	(assert (not (= a 4)))
(a)	(b)	(c)

Fig. 9. SMT problems. (a) SMT problem based on the first trace. (b) SMT problem based on a second trace. (c) SMT problem built from the preceding two problems.

Lemma 1. *For a well-formed machine state m ,*

$$\text{status}(m) = \text{failure} \iff \text{ANS}(\text{getSMT}(m)) = \text{SAT}$$

Proof. The generated SMT problem in Figure 7 and Figure 8 captures the execution trace from the initial state up to the final state by looking at the following facts. First, the machine step in Figure 7 adds an HB relation for two consecutive program locations ρ_0 and ρ_1 in an identical thread, with respect to the program order constraint. Second, the assume, assert and assign expression evaluations in Figure 8 add assert constraints for variable x or expression e with respect to the trace constraint. Note that the negated system is used in the assert evaluation when generating the SMT statement. Third, the sndi and rcvi commands in Figure 8 add any_0 as the definition of a receive or send operation, followed by any_1 as the initialization of operation fields. Finally, the wait(rcvi) command in Figure 8 adds a match pair for the receive operation a_r and the send operation a_s where a_s is obtained by applying the process we described in Section 3.2. Furthermore, an HB relation is added with respect to the program order, such that the messages from a common endpoint are non-overtaking. The observations described above either adds program order constraints for the execution trace modeled in the machine states, or adds constraints for variables or expressions that are also modeled in the machine state. In a word, the generated SMT problem captures the execution trace in the machine.

Notice the fact above, we suppose the status of the final machine state is **failure**. From the reduction rules, we know that there exists at least one assert action that is evaluated false. In Figure 8, we know that expression “ e ” is negated and added to the SMT encoding for the second reduction rule. Since e is evaluated false, ($\text{not } e$) is evaluated true trivially in the generated SMT problem. Other statements are all evaluated true due to the trace capture fact that the program order is assigned correctly based on the execution trace modeled in the machine states, and the variables and expressions are also evaluated true in assume and assign commands. Thus, the SMT problem is evaluated **SAT**. On the other hand, suppose the SMT problem is evaluated **SAT**, indicating that all statements of the SMT problem are satisfiable. Since the reduction rules in Figure 8 add statements with respect to the transition flow of the trace, it implies that each transition of the trace is executed correctly except for some assert action. Because of the same reason above, we know that expression e for the assert action is evaluated false. Thus, the status of the final machine state is **failure**. \square

As we discussed early, Figure 9 (a) and (b) are generated from the same program with two different traces. The encoding in Figure 9 (c) combines part (a) and (b) into one, the answer of which implies the non-deterministic behavior of the program. By solving the problem in Figure 9 (c), we implicitly solve two problems in Figure 9 (a) and (b) respectively. The following definition defines the “combination” behavior of two SMT problems.

Definition 3. A combination operator $COMB$ for two SMT problems, represented as three-tuples smt_1 and smt_2 , where $smt_1 = (\text{defs constraints } ML_1)$ and $smt_2 = (\text{defs constraints } ML_2)$, returns a new SMT tuple, such that,

$$COMB(smt_1, smt_2) = (\text{defs constraints } ML_{new})$$

where ML_{new} is a new relation such that $\forall \mu \in \text{dom}(ML_1) \cup \text{dom}(ML_2)$, $ML_{new}(\mu) = ML_1(\mu) \cdot ML_2(\mu)$. Note that two SMT problems are assumed to hold the same *defs* and *constraints* excluding the *match constraints*.

To further finding the correlation between the combined SMT problem and the single SMT problems, we get the following lemma and proof.

Lemma 2. For a set of traces T_n for the same CTP, and a set of SMT problems $SMT_n = \{smt_0, smt_1, \dots, smt_n\}$, where each member in SMT_n encodes a trace in T_n according to our trace machine, there exists a new SMT problem smt_{total} , where

$$smt_{total} = COMB(smt_n, COMB(smt_{n-1}, COMB(smt_{n-2}, COMB(\dots))))$$

and

$$\text{ANS}(smt_{total}) = \begin{cases} \mathbf{SAT} & \text{iff } \exists smt_i \in SMT_n, \text{ s.t. } \text{ANS}(smt_i) = \mathbf{SAT} \\ \mathbf{UNSAT} & \text{otherwise} \end{cases}$$

Proof. We Prove it by induction.

We consider the base case as a SMT problem that encodes a single trace and the answer can be trivially proved by definition.

Induction. Assume we have combined n SMT problems and the combined SMT problem smt_{total} is evaluated **UNSAT**. We combine an additional SMT problem smt_{n+1} , which is different from any existent SMT problems. If $\text{ANS}(smt_{n+1})$ is equal to **SAT**, the newly combined SMT problem, smt'_{total} , is evaluated **SAT** because the match pairs defined in smt_{n+1} are combined into smt'_{total} . If $\text{ANS}(smt_{n+1})$ is equal to **UNSAT**, on the other hand, smt'_{total} is then evaluated **UNSAT** because all the traces implied in smt'_{total} are evaluated **UNSAT**. Additionally, it can be trivially proved if smt_{total} is evaluated **SAT** then the newly added SMT problem smt_{n+1} do not change the answer. \square

From **Lemma 2**, the combined SMT problem is more powerful because it can find the violation of an assertion among several trace encodings. The following theorem states the relation between the ability of finding violation and the content of the match-list ML .

Theorem 1. *For two SMT problems, $smt_\phi = (\text{defs constraints } ML_\phi)$ and $smt = (\text{defs constraints } ML)$, if $\text{range}(ML_\phi) \subseteq \text{range}(ML)$,*

$$\text{ANS}(smt_\phi) \geq \text{ANS}(smt)$$

Proof. Since the range of ML_ϕ is the subset of the range of ML , we can obtain smt by combining smt_ϕ with other SMT problems. Suppose $\text{ANS}(smt_\phi)$ is equal to **SAT**, by **Lemma 2**, $\text{ANS}(smt)$ is equal to **SAT** as well. If $\text{ANS}(smt_\phi)$ is equal to **UNSAT**, $\text{ANS}(smt)$ is equal to either **UNSAT** or **SAT**, depending on the answers of other single SMT problems that combine smt . In either case, $\text{ANS}(smt_\phi) \geq \text{ANS}(smt)$.

5 Result

To evaluate the reliability and efficiency of our system, we compare our encoding with the work presented in [2]. The encoding in [2] can handle most regular scenarios of CTP and some errors can be found if such traces are modeled. However, It can not always guarantee of finding all hidden errors since the modeling is sometimes incorrect for a specific CTP. By applying the technique in [2] to the program in Figure 1, not both execution traces that we have done in Figure 9 can be modeled. Actually, the encoding in [2] can only model a single trace that matches *snd1* with *rcvA*, followed by matching *send2* with *rcvC*, etc. The main reason for the incorrectly encoding is that the encoding technique in [2] requires *snd1* to happen before *rcvB* and no other candidate send operations for *rcvB* to happen after *snd1*. For our encoding technique, however, we make it flexible that a non-blocking send operation can match a receive action even if some other sends happen before the receive action. As a result, our encoding can correctly model non-deterministic behavior of a concurrent program.

CTP	Example in Figure 1	small1	small2	small3
Our Encoding	17	8	4	11
Encoding in [2]	47	33	18	44

Table 1. Comparison of two encodings

Despite the correctness of our encoding, our encoding is more efficient by comparing the number of the clauses of the encoding itself. For the program in Figure 1, we encodes 17 clauses for modeling the problem, omitting the definition of variables at the beginning of the encoding. Based on the encoding rules in [2], however, there are 47 clauses for the same program, omitting the definitions as well. Other than the program in Figure 1, we manually generate three programs and provide the experimental result in Table 1. Reducing the clause number means a large improvement of efficiency. The result in Table 1 is impressive because the number of our encoding is only $1/5 \sim 1/3$ of that of the encoding in [2]. Thus, we encode more efficient SMT problem for an identical concurrent program.

6 Over-approximated Match-pair Generation

In the previous sections, we know that a precise set of match pair and a CTP is required to verify all possible execution traces through a SMT problem. Generating a precise set of match pair is time consuming because an execution tree needs to be generated, which requires exponential time complexity to finish. It is impossible to verify a regular MCAPI program because of the state explosion. In this section, we present a method that over-approximates the set of match pair, such that match pairs that can exist in the real execution are all included, and some bogus match pairs that can not exist in the execution are also included and required to prune. We then input this over-approximated set of match pair to a CEGAR loop for further verification. Either error is found during one iteration and the process stops, or the current set of match pair is prune to a more precise one for the next iteration.

Other than simulating all possible executions of the CTP, the algorithm that generates the over-approximated match pairs is basically a strategy of matching and pruning among all possible Sndi and Rcvl commands with matched end points. We use the structure in Figure 10 to support the algorithm. We keep one copy of both lists and store the Rcvl and Sndi commands as follows: The Rcvl commands in all threads are scanned and stored in the list of Rcvl, where the index of a Rcvl command, r_i , is the actual order in the thread where this Rcvl command exists. For example, $Rcvl(t_0, i)$ represents the i_{th} Rcvl command in Thread t_0 . Similarly, the Sndi commands in all threads are stored and ordered in the list of Sndi, where s_i is the index of a Sndi command in the sublist of Sndi commands with same from-end-point and to-end-point. For example,

```

([0 Rcv_i_00, Rcv_i_01, ...]
 [1 Rcv_i_10, Rcv_i_11, ...]
 ...
 [n Rcv_i_n0, Rcv_i_n1, ...])

```

(a)

```

((0 [1 Send_i_100, Send_i_101, ...], [2 Send_i_200, Send_i_201, ...], ...))
(1 [0 Send_i_010, Send_i_011, ...], [2 Send_i_210, Send_i_211, ...], ...)
...
(n [0 Send_i_0n0, Send_i_0n1, ...], [1 Send_i_1n0, Send_i_1n1, ...], ...))

```

(b)

Fig. 10. Structure of pending lists. (a) List of Rcv_i commands. (b) List of Snd_i commands.

$Snd_i(t_0, t_1, j)$ represents the j_{th} Snd_i command sending from t_1 to t_0 . In addition, we have the following lemma that supports the pruning strategy in the algorithm.

Lemma 3. *For each Rcv_i command with index r_i and all matched Snd_i commands of this Rcv_i command, if s_j is the index of any Snd_i commands,*

1. $r_i \geq s_j$;
2. $r_i \leq s_j + (N - sendsize_j)$.

Where N is the number of Snd_i commands in all threads, and $sendsize_j$ is the size of the sublist where this Snd_i command exists.

The algorithm that generates the over-approximated set of match pairs works as follows:

- 1 For each Rcv_i command $Rcv_i(t_s, i)$, all Snd_i commands in the entire set of sublists $Snd_iList(t_s)$ are selected as candidates for matching the Rcv_i command.
- 2 For each candidate Snd_i command $Snd_i(t_s, t_d, j)$, check if Lemma 3 is satisfied. If it is satisfied, the pair $(Rcv_i(t_s, i), Snd_i(t_s, t_d, j))$ is added to the result set of match pairs. Otherwise, this pair is ignored, and other candidate Snd_i commands for $Rcv_i(t_s, i)$ are checked.
- 3 Repeat step 1 and 2 until each Rcv_i command is checked with all its candidate Snd_i commands.

The generated set of match pairs is an over-approximation of the precise set because it contains all legal match pairs. Furthermore, the pruning strategy restricts the size of the set because some match pairs that can illegally exist are removed.

Now, we analyze the time complexity of this algorithm, which requires several steps to complete. Organizing the contents of pending lists requires linearly scanning the CTP so that the time complexity for this step is $O(N)$. Matching the possible Snd_i commands with the given Rcv_i commands requires $N_0^2 + N_1^2 +$

$\dots + N_n^2 < N^2$ steps, where $N_0 + N_1 + \dots + N_n = N$, so this step requires $O(N^2)$ time complexity. Thus, the whole process requires $O(N^2)$ time complexity totally.

7 Related Work

Chao Wang et al. provided a symbolic algorithm that detects errors in all feasible permutations of statements in an execution trace [4]. They partition a concurrent program into a set of concurrent trace program (CTP), which we used in our work. Sharma et al. provided a verification tool, MCC, for verifying the MCAPI programs, but it was later discovered that this method does not completely explore the entire execution space of certain kinds of MCAPI programs [3]. Elwakil et al. created a similar work with ours [1, 2]. We compared two works in the previous section, where their encoding does not always model all feasible execution traces and it lacks the efficiency for solving the SMT problem.

8 Conclusions

We have presented a trace language that can describe a CTP with a single trace on the same CTP. By executing a CTP with the language syntax, it can enter the final state that indicates whether the execution is correct or not. Also, we have defined an SMT problem that encodes the specific CTP, in which the match pair is defined to describe the transition of send and receive actions. Then we have proved the consistency of a trace language semantic and an SMT problem for a single trace given the same CTP. By expanding the set of match pair, we have proved that SMT problem can further imply more execution traces for a CTP. Thus, given a precise set of match pair, we can model all possible traces through an SMT problem. Finally, we provide a method of over-approximating the set of match pair that can lead to a refinement loop for finding errors.

9 Future Work

The generation of set of match pair is essential to the efficiency of our refinement process. Currently we have an over-approximated generation method, but the efficiency is not impressive. New methods for generating a much more precise set of match pair are required.

References

1. Elwakil, M., Yang, Z.: CRI: Symbolic debugger for mcaapi applications. In: Automated Technology for Verification and Analysis (2010)
2. Elwakil, M., Yang, Z.: Debugging support tool for mcaapi applications. In: PADTAD '10: Proceedings of the 8th Workshop on Parallel and Distributed Systems (2010)

3. Sharma, S., Gopalakrishanan, G., Mercer, E., Holt, J.: Mcc - a runtime verification tool for mcapi user applications. In: FMCAD (2009)
4. Wang, C., Chaudhuri, S., Gupta, A., Yang, Y.: Symbolic pruning of concurrent program executions. In: ESEC/FSE. pp. 23–32. ACM, New York, NY, USA (2009)