

Proving MCAPI Executions are Correct

Applying SMT Technology to Message Passing

Yu Huang, Eric Mercer, and Jay McCarthy *

Brigham Young University
{yuHuang, egm, jay}@byu.edu

Abstract. Asynchronous message passing for C language is an important paradigm in writing applications for embedded heterogeneous multicore systems. The Multicore Association (MCA), an industry consortium promoting multicore technology, is working to standardize message passing into a single API, MCAPI, for bare metal implementation and portability across platforms. Correctness in such an API is difficult to reason about manually, and testing against reference solutions is equally difficult as reference solutions implement an unknown set of allowed behaviors, and programmers have no way to directly control API internals to expose or reproduce errors. This paper encodes an MCAPI execution as an SMT problem, which if satisfiable, yields a feasible execution schedule on the same trace, such that it resolves non-determinism in the MCAPI runtime in the way that it now fails user provided assertions. Such an encoding is useful for test, debug, and verification of MCAPI program execution. The novelty in the encoding is the use of match pairs (potential send and receive couplings) to model the MCAPI execution in the SMT problem. Match-pair encoding for MCAPI executions, when compared to other encoding strategies, is simpler to reason about, results in significantly fewer terms in the SMT problem, and captures feasible behaviors that are ignored in previously published techniques.

Keywords: Abstraction, refinement, SMT, message passing

1 Introduction

Embedded devices fill all sorts of crucial roles in our lives. They exist as medical devices, as network infrastructure, and they control our automobiles. Embedded devices continue to become more powerful as computing hardware becomes smaller and more modular. It is now becoming commonplace to find multiple processing units inside a single device. The Multicore Association (MCA) [2] is an industry group that has formed to define specifications for low-level communication, resource management, and task management for multicore devices.

One specification that the MCA has released is the Multicore Association Communications API (MCAPI) [3]. The specification defines basic data type,

* Special thanks to Christopher Fischer, formally at BYU and now at Amazon

data structures, and functions that can be used to perform simple message passing operations between different computing entities within a device. Messages can be passed across persistent channels that force an ordering of the messages, or they can be passed to specific *endpoints* within the system. The specification places few ordering constraints on messages passed from one endpoint to another. This freedom introduces into the system possibilities that two or more message are racing if the order of their arrival at the destinations are non-deterministic [14]. Without a way to explore this non-determinism in the MCAPI runtime, it is not possible to test and debug the programs.

Sharma et al. created a method of using concrete execution to verify MCAPI programs, but it was later discovered that this method does not completely explore the entire execution space of certain kinds of MCAPI programs [15]. The method provides match pairs – couplings for potential sends and receives that we use in our work. Instead of exploring all relevant interleavings of an MCAPI program in the concrete execution in [15], Wang et al. provided a symbolic algorithm that verifies each partitioned concurrent trace program with shared memory semantics (CTP) using a satisfiability solver [17]. Elwakil et al. also defined a method of representing MCAPI program executions as satisfiability problems building on the method of [17] adapting it to message passing, but this method does not correctly model all kinds of MCAPI program executions [7]. A shortcoming of this method is the assumption that the user is able to provide the exact set of match pairs. Such an assumption is not reasonable for large complex program traces. In this paper, however, we assume a user provides a possible set of match pairs as input to an MCAPI program. As such, the existing technique either faces the state explosion problem [15], or provides incorrect encoding that does not capture the true program behavior in an MCAPI runtime [7].

We present an encoding for MCAPI program executions that uses match pairs from [15] in a new Satisfiability Modulo Theory (SMT) encoding that requires fewer terms than [7] but includes all possible program executions unlike [7]. In other words, the solution for our encoding provides efficiency. Most importantly, this method correctly captures all possible execution traces allowed by the MCAPI specification in a trace of a concurrent program as an SMT problem enabling the exploration of inherent non-determinism that may exist in any MCAPI runtime implementation. Our main contributions in this paper include:

1. A correct and efficient SMT encoding of an MCAPI program execution; and
2. An $O(n^2)$ algorithm to generate an over-approximation of possible match-pairs.

The rest of the paper is organized as follows: Section 2 presents an MCAPI program in which two program traces exist due to non-determinism in the MCAPI runtime. Section 3 defines an SMT encoding for an MCAPI program execution that captures non-determinism in message order by defining the set of possible match pairs between sends and receives on a common endpoint. Section 4 solves the outstanding problem in other encodings of generating feasible match-pairs to use in the encoding. It presents a $O(n^2)$ algorithm that over-approximates the true match set. It then uses our encoding in a framework that

provides a *CEGAR*¹ loop to prune infeasible match pairs from the initial set until we witness a true error or conclude no error is possible. Section 5 presents the experimental results that show our encoding to be correct and efficient, as other existing encodings both omit critical program behaviors and require more SMT clauses. Section 6 discusses related work, and Section 7 presents our conclusions and future work.

2 Example

It is a challenge to explain intended behavior in simple scenarios consisting of a handful of calls when dealing with concurrency. Consider the MCAPI program in Figure 1 that includes three tasks that use send (`mcapi_msg_send_i`) and receive (`mcapi_msg_rcv_i`) calls to communicate with each other. Line numbers appear in the first column for each task, and the declarations of the local variables are omitted for space. Picking up the scenario just after the endpoints are defined, lines 02 and 05 of task 0 receive two messages on the endpoint *e0* in variables *A* and *B* which are converted to integer values and stored in variables *a* and *b* on lines 04 and 07; task 1 receives one message on endpoint *e1* in variable *C* on line 03 and then sends the message “1” on line 05 to *e0*; and finally, task 2 sends messages “4” and “Go” on lines 04 and 06 to endpoints *e0* and *e1* respectively. Task 0 has additional code (lines 08 - 09) to assert properties of the values in *a* and *b*. The `mcapi_wait` calls block until the associated send or receive buffer is able to be used. Given the scenario, we might ask the question: “*What are the possible values of a and b after the scenario completes?*”

The intuitive trace is presented in the first three columns of Figure 2. Note that the first column contains the task number in Figure 1. The second column presents the commands identified by the line numbers shown in Figure 1. Also, we define a shorthand for each command of send (denoted as S), receive (denoted as R), or wait (denoted as W) in the third column for convenience in the presentation. For each command $O_{i,j}(k, \&h)$, $O \in \{S, R, W\}$, *i* represents the task number, *j* represents the command line, *k* represents the destination task number, and *h* represents the command handler. Note that a specific destination task number is assigned to each receive, which implies the matched send in the program runtime. To provide the program non-determinism, the destination task number can be replaced with “*” in an MCAPI program. From the trace, variable *a* should contain 4 and variable *b* should contain 1 since task 2 must first send message “4” to *e0* before it can send message “Go” to *e1*; consequently, task 1 is then able to send message “1” to *e0*. The assume notation asserts the control flow taken by the program execution which in this example, asserts the true branch of the condition on line 08 of task 0. At the end of execution the assertion on line 09 of task 0 holds and no error is found. Such intuition is a valid program execution.

To complete a send call, the message can be either copied out into a runtime provided buffer or to an endpoint provided buffer [16]. We call the first(second)

¹ The *CEGAR* loop is defined as a Counter Example Guided Abstraction Refinement.

message buffer *infinite-buffer(zero-buffer)*. It is remarkable that *infinite-buffer* provides more non-deterministic behaviors, because a particular endpoint can receive more messages under the circumstance of *infinite-buffer* while the *zero-buffer* can only allow one message to be received by a specific endpoint at a time. For our example in Figure 1, if we use *zero-buffer* as the buffer setting, the send call on line 05 of task 1 can not match with the receive call on line 02 of task 0. This is because the send call on line 04 of task 2 must be completed before the send call on line 06 of task 2 can match with the receive call on line 03 of task 1. When using *infinite-buffer*, however, we can have another scenario in the fourth and fifth column of Figure 2 written in the shorthand notation. The variable *a* contains 1 instead of 4, since the message “1” is sent to *e0* after sending the message “Go” to *e1* as it is possible for the send on line 04 of task 2 to be delayed in transit. The specification indicates that the wait on line 05 of task 2 returns once the buffer is available. That only means the message is somewhere in the MCAPI runtime; it does not mean the message is delivered. As such, it is possible for the message to be delayed in transit allowing the send from task 1 on line 05 to arrive at *e0* first and be received in variable “*a*”. Such a scenario is a program execution that results in an assertion failure at line 09 in Figure 1. From the discussion above, it is important to consider non-determinism in the MCAPI runtime when testing or debugging an MCAPI program execution. Note that we use *infinite-buffer* for the program execution in the rest of our discussion. The next section presents our method of SMT encoding that follow the same control flow path through an MCAPI program. The encoding can be solved by an SMT solver such as Yices [5] and Z3 [13], and the non-deterministic behavior of the program can be resolved.

Task 0	Task 1	Task 2
<pre> 00 initialize(NODE_0,&v,&s); 01 e0 = create_endpoint(PORT_0,&s); 02 msg_rcv_i(e0,A,sizeof(A),&h1,&s); 03 wait(&h1,&size,&s,MCAPI_INF); 04 a = atoi(A); 05 msg_rcv_i(e0,B,sizeof(B),&h2,&s); 06 wait(&h2,&size,&s,MCAPI_INF); 07 b = atoi(B); 08 if(b > 0); 09 assert(a == 4); 10 finalize(&s); </pre>	<pre> 00 initialize(NODE_1,&v,&s); 01 e1 = create_endpoint(PORT_1,&s); 02 e0 = get_endpoint(NODE_0,PORT_0,&s); 03 msg_rcv_i(e1,C,sizeof(C),&h3,&s); 04 wait(&h3,&size,&s,MCAPI_INF); 05 msg_send_i(e1,e0,"1",2,N,&h4,&s); 06 wait(&h4,&size,&s,MCAPI_INF); 07 finalize(&s); </pre>	<pre> 00 initialize(NODE_2,&v,&s); 01 t2 = create_endpoint(PORT_2,&s); 02 t0 = get_endpoint(NODE_0,PORT_0,&s); 03 t1 = get_endpoint(NODE_1,PORT_1,&s); 04 msg_send_i(e2,e0,"4",2,N,&h5,&s); 05 wait(&h5,&size,&s,MCAPI_INF); 06 msg_send_i(e2,e1,"Go",3,N,&h6,&s); 07 wait(&h6,&size,&s,MCAPI_INF); 08 finalize(&s); </pre>

Fig. 1. An MCAPI concurrent program

3 SMT Model

The novelty of the SMT encoding in this paper is its use of match pairs rather than the state-based or order-based encoding of prior work [7, 6]. The algorithm

Trace 1			Trace 2		
Task	Command	Shorthand	Task	Command	
2	04 msg_send_i(e2, e0, "4", 2, N, &h5, &s);	S _{2,4} (0, &h5)	2	04 S _{2,4} (0, &h5)	
2	05 wait(&h5, &size, &s, MCAPI_INF);	W(&h5)	2	05 W(&h5)	
0	02 msg_rcv_i(e0, A, sizeof(A), &h1, &s);	R _{0,2} (2, &h1)	2	06 S _{2,6} (1, &h6)	
0	03 wait(&h1, &size, &s, MCAPI_INF);	W(&h1)	2	07 W(&h6)	
2	06 msg_send_i(e2, e1, "Go", 3, N, &h6, &s);	S _{2,6} (1, &h6)	1	03 R _{1,3} (2, &h3)	
2	07 wait(&h6, &size, &s, MCAPI_INF);	W(&h6)	1	04 W(&h3)	
0	04 a = atoi(A);		1	05 S _{1,5} (0, &h4)	
1	03 msg_rcv_i(e1, C, sizeof(C), &h3, &s);	R _{1,3} (2, &h3)	1	06 W(&h4)	
1	04 wait(&h3, &size, &s, MCAPI_INF);	W(&h3)	0	02 R _{0,2} (1, &h1)	
1	05 msg_send_i(e1, e0, "1", 2, N, &h4, &s);	S _{1,5} (0, &h4)	0	03 W(&h1)	
1	06 wait(&h4, &size, &s, MCAPI_INF);	W(&h4)	0	04 a = atoi(A);	
0	05 msg_rcv_i(e0, B, sizeof(B), &h2, &s);	R _{0,5} (1, &h2)	0	05 R _{0,5} (2, &h2)	
0	06 wait(&h2, &size, &s, MCAPI_INF);	W(&h2)	0	06 W(&h2)	
0	07 b = atoi(B);		0	07 b = atoi(B);	
0	08 assume(b > 0);		0	08 assume(b > 0);	
0	09 assert(a == 4);		0	09 assert(a == 4);	

Fig. 2. Two execution traces of the MCAPI program in Figure 1

to create the encoding takes as input a set of possible match pairs and a trace through an MCAPI program with the appropriate assumes and asserts as shown in Figure 2. A match pair is the coupling of a receive to a particular send. Figure 3(a) is the set of possible match pairs for the program in Figure 1 using our shorthand notation defined in Figure 2. The set admits, for example, that $R_{0,2}$ can be matched with either $S_{1,5}$ or $S_{2,4}$. The SMT encoding in this paper asks the SMT solver to resolve the match pairs for the system in such a way that the final values of program variables meet the assumption on control flow but violate some assertion.

Before presenting the SMT encoding for a specific example, we need to explain auxiliary data structures and support functions in detail. We give the definitions in the Yices input language.

$$\begin{aligned}
 &(\text{define HB} :: (\rightarrow a::\text{int } b::\text{int} \\
 &\quad (\text{subtype } (c::\text{bool}) (\text{ite } (< a b) (= c \text{ true}) (= c \text{ false})))) \quad (1)
 \end{aligned}$$

The HB function in equation (1) creates a happens-before relationship between two events. The notation “ite” refers to a condition statement such that if the condition is satisfied, then the first subclause is true; otherwise, the second subclause is true. The function takes two events a and b as parameters, such that $a, b \in \mathbb{N}$, where \mathbb{N} is the set of natural numbers. In the SMT encoding, every location is assigned an event, and HB orders those locations (i.e., program order, etc.). This function creates a constraint such that the first event must be less than the second event, indicating that the first event occurs before the second event. For example, action a happens before b if and only if $event_a$ for a is greater than $event_b$ for b . This function is not only used to assert the program order of statements within the same thread, but also to assert, for any matched pair, that the send operation occurs before the receive.

$$\begin{aligned}
&(\text{define-type } R \text{ (record } M::\text{int } e::\text{int } value::\text{int } event::\text{int})) \\
&(\text{define-type } S \text{ (record } O::\text{int } ID::\text{int } e::\text{int } value::\text{int } event::\text{int}))
\end{aligned} \tag{2}$$

We define two records in equation (2) that are essential to encoding the SMT problem. The first record R defined above represents a receive action such that $M, e, value, event \in \mathbb{N}$. M represents a free variable that is assigned the $event$ field of the matched record S . The endpoint e is the endpoint of the receive operation, $value$ represents a value received by the receive operation, and $event$ is the event term which indicates the program order for the receive action associated with the record. The second record S defined above is a send action such that $O, ID, e, value, event \in \mathbb{N}$. The fields e and $event$, as those in R , are the destination endpoint and event term, respectively. The field $value$ is the value being sent. O is a free variable used for match pairs, and it is assigned an event term of a receive operation if the receive operation is to be matched to the send operation. ID is a unique number that identifies each send operation. Note that the integer value assigned to $event$ indicates the program order. Also, we use two match variables, M and O , for different purposes. M is used for preventing two receive operations from being matched with the same send operation. O is used for ordering two receive operations with respect to the program order of two matched send operations. The SMT solver must assign the M and $event$ fields of the send tuple, and it must assign the $event$, M , and $value$ fields of the receive tuple in a way that is consistent with all the constraints in the encoding.

$$\begin{aligned}
&(\text{define MATCH} :: (\rightarrow r::R s::S (\text{subtype } (c::\text{bool}) (\text{ite} \\
&\quad (\text{and } (= (\text{select } r \text{ } e) (\text{select } s \text{ } e)) \\
&\quad \quad (= (\text{select } r \text{ } value) (\text{select } s \text{ } value)) \\
&\quad \quad (\text{HB } (\text{select } s \text{ } event) (\text{select } r \text{ } event)) \\
&\quad \quad (= (\text{select } r \text{ } M) (\text{select } s \text{ } ID)) \\
&\quad \quad (= (\text{select } s \text{ } O) (\text{select } r \text{ } event)))) \\
&\quad (= c \text{ true}) \\
&\quad (= c \text{ false}))))
\end{aligned} \tag{3}$$

The MATCH function in equation (3) takes tuple R and S as parameters. Each “select” statement selects a field of the tuple R or S . The function creates a number of constraints that represent the pairing of the specified send and receive operations. This function asserts that the destination endpoint of the send operation is the same as the endpoint used by the receive operation, that the value sent by the send operation is the same value received by the receive operation, that the send operation occurs before the receive operation, that the “match” value in the receive tuple is equal to the “ID” value in the send tuple, and that the “O” value in the send tuple is equal to the “event” value in the receive tuple. Note that the encoding is going to use the “O” value in conjunction with the HB function to enforce message non-overtaking when messages are sent from the same endpoint. Consider a simple example below that sends two

messages from a Task 0 to Task 1,

Task 0	Task 1
$S_{0,1}(1, \&h1)$	$R_{1,1}(*, \&h3)$
$S_{0,2}(1, \&h2)$	$R_{1,2}(*, \&h4)$
$W(\&h1)$	$W(\&h3)$
$W(\&h2)$	$W(\&h4)$

The “O” values in the send records will be assigned to the order tracking events for $R_{1,1}$ and $R_{1,2}$ in the final encoding. Message non-overtaking from transactions on common endpoints is encoded by forcing the send events to be received in program order using the HB function as below in our simple example:

$$(\text{HB } (\text{select } S_{0,1} \text{ O}) (\text{select } S_{0,2} \text{ O}))$$

$$\begin{aligned}
&(\text{define NE} :: (\rightarrow r1::R \ r2::R \ (\text{subtype } (c::\text{bool}) \ (\text{ite} \\
& \quad (= (\text{select } r1 \ M) (\text{select } r2 \ M)) \\
& \quad (= c \ \text{false}) \\
& \quad (= c \ \text{true}))))))
\end{aligned} \tag{4}$$

The NE function in equation (4) is used to assert that no two receive operations are matched with the same send operation. The parameters for this function are two receive tuples. An assertion is made that the values of their “match” fields, M , are not equal. This shows that they are paired with two different send operations.

Prior SMT models of MCAPI program executions implicitly compute match pairs by adding possible happens-before relations on sends and receives with the conditions under which those orderings are valid. The SMT solver is then asked to resolve the happens-before relation by choosing specific orders of sends and receives. Match pair encoding, as presented in this work, though it has the same computational complexity as order based encoding (you still need to figure out match pairs on endpoints), is simpler to reason about directly rather than implicitly through orders, results in significantly fewer terms in the SMT problem, and does not restrict out possible matches allowed by the MCAPI specification of the runtime as do existing encodings.

The presentation of the SMT encoding is structured as

$$smt = (\text{defs } constraints \ match)$$

where the field *defs* represents all definitions of the send, receive operations and the free variables; The field *constraints* represents all constraint clauses, including the assert clauses and the clauses built by the HB functions; and the field *match* represents the clauses built by the MATCH function on a set of match pairs and the NE function.

The *Intra-Happens-Before-Order(Intra-HB)* in [16] constrains the Happens-Before relations between commands in any control flow path. Those relations are essential to the POE algorithm for their Message-passing verifier. Our SMT

structure encodes all the relations in *Intra-HB*, such that it can follow any control flow path executed by the POE algorithm in [16]. The *Intra-HB* consists of 8 *Happens-Before* relations:

$$\prec_{mhb} = \prec_S \cup \prec_R \cup \prec_W \cup \prec_{SW} \cup \prec_{RW} \cup \prec_{mb} \cup \prec_{fo} \cup \prec_{io}$$

The relations for *Barrier* commands are omitted because they are not supported in the MCAPI programs. The relations \prec_S , \prec_R and \prec_W are used to constrain the state transition of commands. For example, issuing a send command happens before matching it with some receive in transit. Instead of explicitly giving the state information, the SMT structure in this paper implicitly encodes the state transitions in the field *constraints*. The relations \prec_{SW} , \prec_{RW} , \prec_{fo} and \prec_{io} constrain the program order of commands. We encode those relations in the field *constraints* of our SMT encoding. The relation \prec_{mb} builds the match before for sends and receives. Also, we have the relation in the field *match* in our SMT encoding.

The input for generating the SMT encoding is a MCAPI program together with a set of match pairs. The program in Figure 1 with the set of match pairs in Figure 3(a) is such an example. Given an input, the algorithm for generating our encoding is enumerated as follows.

1. Define a variable for the program order of each assume or assert command in the field *defs*;
2. for each assume, add a statement to *constraints*, and for each assert, add a negated assert to *constraints* in order to obtain an execution trace with assertion failure from an SMT solver;
3. for each send or receive in the trace, create an *S* or *R* definition and create an associated “.event” variable for program order and add these to *defs*;
4. for each program order represented by the field event in each thread, add an HB relation on the associated event for program order to *constraints*;
5. for each match pair for a send and receive as input, add an MATCH relation to *match*;
6. for two sends $S_{i,j}$ and $S_{i,k}$ in the same task i sent to an identical endpoint such that $i < k$, if the relation (HB $S_{i,j}.\text{event } S_{i,k}.\text{event}$) is defined in *constraints*, add the clause (HB (select $S_{i,j}$ O) (select $S_{i,k}$ O)) to *constraints*.
7. collect into a disjunction for all possible send matches for any given receive in the field *match*;
8. and collect into a NE relation for all receives that are matched with any identical send and add it to *match*.

Figure 3(b) is a concrete example of the SMT encoding for our running program in Figure 1 given the set of match pairs in Figure 3(a) as input. The field *defs* is not shown because the definitions are not novel to our solution. Line 00 - 12 are the *constraints*, and line 13 - 18 are the *match* of the SMT encoding. In particular, by encoding function HB, $R_{0,2}$ happens before the wait command $W(\&h1)$ given the “event” for each command as parameters on line 00 in Figure 3(b). We use the wait command with handler in Figure 2 to denote

	<i>defs</i> is not shown;
	<i>constraints</i> ;
	00 (HB R02.event W(&h1).event)
	01 (HB W(&h1).event R05.event)
	02 (HB R05.event W(&h2).event)
	03 (HB W(&h2).event assume.event)
	04 (HB assume.event assert.event)
	05 (HB R13.event W(&h3).event)
(R02 S24)	06 (HB W(&h3).event S15.event)
(R02 S15)	07 (HB S15.event W(&h4).event)
	08 (HB S24.event W(&h5).event)
(R05 S24)	09 (HB W(&h5).event S27.event)
(R05 S15)	10 (HB S27.event W(&h6).event)
	11 (assert (> b 0))
(R13 S27)	12 (assert (not (= a 4)))
	<i>match</i> ;
	13 (or (MATCH R02 S24)
	14 (MATCH R02 S15))
	15 (or (MATCH R05 S24)
	16 (MATCH R05 S15))
	17 (MATCH R13 S27)
	18 (NE R02 R05)

(a)
(b)

Fig. 3. A match pair set and SMT encoding of the system in Figure 1. (a) The match pairs based on endpoints. (b) The SMT encoding where **HB** creates a happens-before constraint, **MATCH** creates a match pair constraint, and **NE** creates a constraint to preclude the use of conflicting match pairs.

the wait command for a send or receive command. Any assignment to event variables by the SMT solver must comply with program order constraints. Each event in the HB orderings on line 00 - 10 are ordered in a thread of Figure 1. Also, the **MATCH** function encodes the relation such that $R_{0,2}$ is matched with $S_{1,5}$ on line 14 in Figure 3(b). The **NE** term precludes $R_{0,2}$ and $R_{0,5}$ from matching to the same send on line 18. Other than the support functions, the first assert on line 11 works as the assume (line 08 at task 0) in Figure 1, such that it prevents the SMT solver from finding solutions that are not consistent with control flow which requires “ $b \geq 0$ ”. Finally, the second assert on line 12 is negated as the goal is to find schedules that violate the property. In our example, the interest is in solutions that resolve in $a \neq 4$ and follow the same control flow with $b > 0$.

Other than the basic structure of the SMT encoding, we use the function

$$\text{ANS}(smt) \mapsto \{\mathbf{SAT}, \mathbf{UNSAT}\}$$

to return the solution of an SMT problem, such that **SAT** represents a satisfiable solution that finds an execution of the MCAPI program that violates the user defined correctness property, and **UNSAT** represents an unsatisfiable solution that all possible execution traces either meet correctness property in the same control flow, or follow a different control flow. Note that **UNSAT** and **SAT** are ordered such that **UNSAT** < **SAT**.

The SMT encoding defined above is used to capture the non-deterministic behavior of an MCAPI program by giving a complete set of match pairs. As we discussed in the previous section, the MCAPI program in Figure 1 contains two

outcomes of execution. The SMT encoding we present in Figure 3(b) captures both executions, since the set of match pairs in Figure 3(a) is a complete set where all matches that can occur in the real execution for our running example in Figure 1 are included, and all matches that can not occur in the real execution are not included. Furthermore, the following theorem states that we can over-approximate the true set of match pairs and still prove correctness. If there is no error with the over-approximated set, then there is no error arising from non-determinism in the runtime on that program execution. If there is an error from the over-approximated set, then that error may or may not be a real error. Note that two SMT problems smt_α and smt_β in the following theorem have an identical *defs* and *constraints* sets, and the match set of smt_α is the subset of that of smt_β so that smt_β represents an over-approximation of smt_α .

Theorem 1. *For two SMT problems, $smt_\alpha = (defs\ constraints\ match_\alpha)$ and $smt_\beta = (defs\ constraints\ match_\beta)$, if $set(match_\alpha) \subseteq set(match_\beta)$,*

$$ANS(smt_\alpha) \leq ANS(smt_\beta)$$

where $set(match)$ represents the input set of match pairs for the **MATCH** clauses in the field *match*.

<pre> <i>defs</i> is not shown; <i>constraints</i>; 00 (HB R02.event W(&h1).event) 01 (HB W(&h1).event R05.event) 02 (HB R05.event W(&h2).event) 03 (HB W(&h2).event assume.event) 04 (HB assume.event assert.event) 05 (HB R13.event W(&h3).event) 06 (HB W(&h3).event S15.event) 07 (HB S15.event W(&h4).event) 08 (HB S24.event W(&h5).event) 09 (HB W(&h5).event S27.event) 10 (HB S27.event W(&h6).event) 11 (assert (> b 0)) 12 (assert (not (= a 4))) <i>match</i>; 13 (MATCH R02 S24) 14 (MATCH R05 S15) 15 (MATCH R13 S27) </pre>	<pre> <i>defs</i> is not shown; <i>constraints</i>; 00 (HB R02.event W(&h1).event) 01 (HB W(&h1).event R05.event) 02 (HB R05.event W(&h2).event) 03 (HB W(&h2).event assume.event) 04 (HB assume.event assert.event) 05 (HB R13.event W(&h3).event) 06 (HB W(&h3).event S15.event) 07 (HB S15.event W(&h4).event) 08 (HB S24.event W(&h5).event) 09 (HB W(&h5).event S27.event) 10 (HB S27.event W(&h6).event) 11 (assert (> b 0)) 12 (assert (not (= a 4))) <i>match</i>; 13 (MATCH R02 S15) 14 (MATCH R05 S24) 15 (MATCH R13 S27) </pre>
(a)	(b)

Fig. 4. Two SMT encodings of the system in Figure 1. (a) The SMT encoding for Trace 1 in Figure 2. (b) The SMT encoding for Trace 2 in Figure 2.

Proof Sketch Consider the MCAPI program in Figure 1 as an example. Figure 4(a) and (b) are two different SMT encodings for our running example in

Figure 1 generated from different sets of possible match pairs, such that Figure 4(a) encodes trace 1 in Figure 2 and Figure 4(b) encodes trace 2 in Figure 2. By solving the encodings in Figure 4(a) and (b) for trace 1 and 2 in Figure 2 respectively, we get an unsatisfiable solution for Figure 4(a), and a satisfiable solution for Figure 4(b). As we discussed in Section 2, trace 1 is an execution without failure of the assertion, and trace 2 is the one that fails the assertion. Compare both encodings with that in Figure 3(b), we find that the fields *defs* and *constraints* are identical except for the *MATCH* clauses. In particular, the input set of match pairs of either Figure 4(a) or (b) is the subset of that in Figure 3(b). As discussed above, the encoding in Figure 3(b) captures the non-deterministic behavior of our running program in Figure 1, which encodes trace 1 and 2 in Figure 2 into one single SMT problem. Thus, an SMT solver will return a satisfiable solution for the encoding in Figure 3(b). Thus, ANS on the encoding in Figure 3(b) is greater than or equal to the ANS on the encoding in either Figure 4(a) or (b). The formal proof is given in a long version of paper at (<http://students.cs.byu.edu/~yhuang2/downloads/paper.pdf>).

The non-determinism of an MCAPI program discussed above relies on the *infinite-buffer* setting. If the *zero-buffer* setting is used instead, the SMT encoding needs to be reorganized as follows. For each match pair that can not exist in the program runtime caused by the *zero-buffer* problem discussed in Section 2, add extra HB clauses in the SMT encoding to prevent incorrect behavior. For example in Figure 1, to prevent $R_{1,3}$ to be matched with $S_{2,6}$ before $R_{0,2}$ is matched with $S_{2,4}$ prohibited by the *zero-buffer*, an HB relation is added such that the wait command $W(\&h1)$ for $R_{0,2}$ happens before $S_{0,6}$.

4 Generating Match Pairs

The exact set of match pairs can be generated by simulating the program trace and using a depth-first search to enumerate non-determinism arising from concurrent sends and receives. Such an effort, however, solves the very problem we are trying to solve because if you simulate the program trace exploring all non-determinism, then you can also verify all runtime choices for property violations. In this section, we present an algorithm that does not require an exhaustive enumeration of runtime behavior in simulation. Our algorithm over-approximates the match pairs such that match pairs that can exist in the runtime are all included and some bogus match pairs that can not exist in the runtime may or may not be included.

The algorithm generates the over-approximated match pair set by matching each pair of the send and receive commands at common endpoints and then pruning obvious matches that can not exist in any runtime implementation of the specification. Figure 5 is an example program in our shorthand notation to present our algorithm. The sends $S_{1,1}$, $S_{1,3}$ and $S_{2,1}$ have Task 0 as an identical destination endpoint. The send $S_{0,3}$ has Task 1 as the destination endpoint. Algorithm proceeds by first linearly traversing each task of the program storing each receive and send command into two distinct structured lists. Sends are

Task 0	Task 1	Task 2
$R_{0,1}(*, \&h1)$	$S_{1,1}(0, \&h5)$	$S_{2,1}(0, \&h8)$
$R_{0,2}(*, \&h2)$	$R_{1,2}(*, \&h6)$	
$S_{0,3}(1, \&h3)$	$S_{1,3}(0, \&h7)$	
$R_{0,4}(*, \&h4)$		

Fig. 5. Another MCAPi concurrent program

structured as in (6) and receives are structured as in (5).

$$\begin{aligned}
(e_0 &\rightarrow ((0, R_{0,1}), (1, R_{0,2}), \dots)) \\
(e_1 &\rightarrow ((0, R_{1,1}), (1, R_{1,2}), \dots)) \\
&\dots \\
(e_n &\rightarrow ((0, R_{n,1}), (1, R_{n,2}), \dots))
\end{aligned} \tag{5}$$

The list in (5) keeps record of the receive commands uniquely identified by the endpoint e_x , where x is the number of the endpoint. The integer in the first field of each pair indicates the program order of the receive commands within each task on the specified endpoint. Note that the receive command with lower program order should be served first in the program runtime. The receive command in the second field of the list is defined as a record R in equation 2 of the previous section. The notations $R_{0,1}$, $R_{0,2}$, etc. represent the unique identifiers for the receives. Recall that we operate on a program trace so there are no loops and all instances of send and receive operations are uniquely identified.

$$\begin{aligned}
&\text{“dst”} \quad \text{“src”} \quad \text{“src”} \\
(e_0 &\rightarrow ((e_1 \rightarrow ((0, S_{1,1}), (1, S_{1,2}), \dots), (e_2 \rightarrow (\dots), \dots)))) \\
(e_1 &\rightarrow ((e_0 \rightarrow ((0, S_{0,1}), (1, S_{0,2}), \dots), (e_2 \rightarrow (\dots), \dots)))) \\
&\dots \\
(e_n &\rightarrow ((e_0 \rightarrow ((0, S_{0,3}), (1, S_{0,4}), \dots), (e_1 \rightarrow (\dots), \dots))))
\end{aligned} \tag{6}$$

Also, we keep record of the send commands using the second list in (6) that is uniquely identified by the destination endpoint e_x . Each sublist for the destination endpoint is uniquely identified by the source endpoint e_y . Similarly, the integer in the first field of each pair indicates the program order of the send command within the same task from a common source and to a common endpoint. The send command in the second field of the list is defined as a record S in equation 2 of the previous section. The notations $S_{1,1}$, $S_{1,2}$, etc. represent the unique identifiers for the sends.

In our running example in Figure 5, we generate our receive list and send list in equation (7) and (8), respectively.

$$\begin{aligned}
(0 &\rightarrow ((0, R_{0,1}), (1, R_{0,2}), (2, R_{0,4}))) \\
(1 &\rightarrow ((0, R_{1,2})))
\end{aligned} \tag{7}$$

$$\begin{aligned}
(0 &\rightarrow ((1 \rightarrow ((0, S_{1,1}), (1, S_{1,3})), (2 \rightarrow ((0, S_{2,1})))))) \\
(1 &\rightarrow ((0 \rightarrow ((0, S_{0,3}))))))
\end{aligned} \tag{8}$$

Note that $R_{0,1}$ is the first receive operation in endpoint 0, and $S_{2,1}$ is the first send from the source endpoint 2 to the destination endpoint 0.

The second step for our algorithm is to linearly traverse both lists to generate match pairs between send and receive commands. Observe that the sends defined with an identical destination endpoint and an identical source endpoint are processed in a FIFO order according to the MCAPI runtime specification. The same is true of receives on a common endpoint in a common task. The FIFO ordering, sometimes referred to as message non-overtaking, lets us prune match pairs that obviously can not be generated in any valid implementation of the MCAPI runtime.

Let I_r and I_s be the program order indication in the lists of (5) and (6) respectively and similarly R and S be the corresponding actions. Note that I_r and I_s are increased by 1 for two consecutive receives or sends, respectively. Further, N_s is the number of send actions in the program and $n_s(src, dst)$ be the number of sends from src to dst , where src and dst represent the source and destination endpoints, respectively. A send and received can be matched if and only if

1. the destination endpoint of S is identical to the endpoint of R ;
2. $I_r \geq I_s$;
3. and $I_r \leq I_s + (N_s - n_s(src, dst))$.

If all rules of the criteria are satisfied for a send and a receive, we build a match pair for the send and the receive in the result set. In our concrete example, $R_{0,1}$ is matched with $S_{1,1}$ or $S_{2,1}$, but it can not be matched with $S_{1,3}$ since the second rule is not satisfied such that the order for $R_{0,1}$ is less than the order of $S_{1,3}$. We repeatedly apply the criteria to match sends and receives until we completely traverse the lists in (5) and (6). The generated set of match pairs for our example in Figure 5 is imprecise such that some extra match pair, we call “bogus”, that can not exist in the real execution of program are included. In particular, the match pair $(S_{2,1} R_{0,4})$ is a “bogus” match pair because it is not possible to order $S_{1,3}$ before $R_{0,2}$ since $R_{1,2}$ can only match with $S_{0,3}$ that must occur after $R_{0,2}$. Fortunately, the encoding in this paper is strong enough to preclude that bad match pair in any satisfying solution.

Now let us analyze the time complexity of the algorithm. Traversing the tasks linearly takes $O(N)$ to complete, where N is the total lines of code of the program. Traversing the list of receives and the list of sends takes $O(mn)$ to complete, where m is the total number of sends and n is the total number of receives. Note that $m + n \leq N$. Totally, the algorithm takes $O(N + mn) \leq O(N + N^2) \approx O(N^2)$ to complete.

As our match pair generation algorithm over-approximates the true set of match pairs allowed by the runtime for an MCAPI program trace, when used with our SMT encoding it may result in false positives (i.e., errors arising from match pairs that cannot happen in any implementation of the MCAPI specification). To resolve such false positives, the witness trace from any SMT problem that is satisfiable is given to an oracle to validate if the trace is feasible. If it is, then the error is a true error. If it is not feasible in a valid MCAPI runtime, then

the bad match pair is pruned and a new SMT encoding is generated. In this way, the encoding and match-pair generation algorithm can be used in a CE-GAR loop. The oracle to detect infeasible match pairs come from formal model of the MCAPI runtime. The details are found in [12] and long version of our paper at (<http://students.cs.byu.edu/~yhuang2/downloads/paper.pdf>).

5 Results

To evaluate the reliability and efficiency of our system, we compare to [7] even though it misses valid MCAPI runtime behavior because it is the only other encoding for MCAPI programs and it suffices to illustrate the efficiency of our encoding.

	Program Order	Matches	Assume&Assert	Extra Clauses
Our Encoding	11	4	2	0
Encoding in [7]	22	13	3	8

Table 1. Comparison of two encodings for the MCAPI program in Figure 1.

For the program in Figure 1, we encodes 17 clauses for modeling the problem, omitting the definition of variables at the beginning of the encoding. Based on the encoding rules in [7], however, there are 47 clauses for the same program, omitting the definitions as well. Table 1 shows the number of clauses that constraint program order, match pairs, assume/asserts, and any other extra uncategorized clauses. Program order comprises the bulk of the encoding. The extra clauses in [7] come from auxiliary variables in the encoding.

	Example in Figure 1	small1	small2	small3
Our Encoding	17	8	4	11
Encoding in [7]	46	33	18	44

Table 2. Comparison of two encodings for four different MCAPI programs with clause number.

Other than the program in Figure 1, we manually generate three programs and provide the experimental result in Table 2. Small1 is a program with two tasks, where two sends are in the first task, and two receives are in the second task. Small2 is a program with only one task, where one send and one receive are in this task. Small3 is a program with three tasks, where three sends are in the first task, two receives are in the second task, and one receive is in the third task. The number of clauses follows the performance of the SMT solver

with fewer clauses leading to better runtime. Table 2 shows that our encoding generates $1/5 - 1/3$ the clauses of that of the encoding in [7]. We believe these reductions will generalize to any given program trace and match pair set.

6 Related Work

Morse et al. provided a formal modeling paradigm that is callable from C language for the MCAPI interface [12]. This model correctly captures the behavior of the interface and can be applied for model checking analysis for C programs that use the API.

In [16], Vakkalanka et al. provided POE, a partial order algorithm of their dynamic verifier, ISP, for MPI programs. POE explores all interleavings of an MPI program, and is able to detect hidden deadlocks under *zero-buffer* setting. If *infinite-buffer* is applied, however, some interleavings may not be found. In [15], Sharma et al. provided MCC, a dynamic verifier for MCAPI programs. MCC systematically explores all interleavings of an MCAPI program by concretely executing the program repeatedly. MCC uses DPOR [8] to reduce the redundant interleavings of the execution, however, it does not include all possible traces allowed by the MCAPI specification. Wang et al. provided a symbolic algorithm that detects errors in all feasible permutations of statements in an execution trace in the shared memory system [17]. In this method, the program is partitioned into several concurrent trace programs (CTPs), and the encoding for each CTP is verified using a satisfiability solver. Elwakil et al. provided a similar work with ours [6, 7]. In their work, the method of [17] is used and adapted to the message passing system. As shown in the previous section of this paper, their method does not correctly encode all possible execution traces of an MCAPI program.

Other than the application in message passing system, the SMT-based verification can be applied to other systems and programs. [11] presented a precise verification of heap-manipulating programs using SMT solvers. [10] presented some challenges in SMT-based verification for sequential system code, and tackled these issues by extending the standard SMT solvers.

The application of static analysis is another interesting avenue of research to test or debug message passing programs. [18] and [4] are the approaches for MPI [1], another message passing interface standard. [9] presented a system that uses static analysis to determine offline the topology of the communications and nodes in the input MCAPI program.

7 Conclusions and Future Work

We have presented an SMT encoding of an MCAPI program that uses match pairs rather than the state-based or order-based encoding in the prior work. Unlike the existing method of SMT encoding, our encoding can correctly capture the non-deterministic behavior of an MCAPI program. Also, we have provided an algorithm with $O(N^2)$ time complexity that over-approximates the true set of match pairs, where N is the total number of code lines of the program. Once the

match pairs are generated by the algorithm, our solution with a *CEGAR* loop is provided to prune infeasible match pairs from the initial set until we witness a true error or conclude no error is possible. The experimental results shows that our encoding can be solved efficiently compared with the existing work in [7].

The generation of the set of match pairs is essential to the efficiency of our *CEGAR* loop. Currently we have an over-approximated generation method, but the efficiency is not impressive. New methods for generating a much more precise set of match pair are required. Also The mehtod of DPOR are considered to combine with our SMT encoding to generate the execution traces. We can then use the SMT encoding to further improve the DPOR technique.

References

1. MPI: A message-passing interface standard. <http://www.mpi-forum.org/docs/mpi-2.2/index.htm>
2. The multicore association. <http://www.multicore-association.org>
3. The multicore association resource management API. <http://www.multicore-association.org/workgroup/mcapi.php>
4. Bronevetsky, G.: Communication-sensitive static dataflow for parallel message passing applications. In: CGO. IEEE Computer Society (2009)
5. Dutertre, B., de Moura Leonardo: A fast linear-arithmetic solver for DPLL(T). In: CAV. vol. 4144 of LNCS, pp. 81–94. Springer-Verlag (2006)
6. Elwakil, M., Yang, Z.: CRI: Symbolic debugger for mcapi applications. In: Automated Technology for Verification and Analysis (2010)
7. Elwakil, M., Yang, Z.: Debugging support tool for mcapi applications. In: PADTAD '10: Proceedings of the 8th Workshop on Parallel and Distributed Systems (2010)
8. Flanagan, C., Godefroid, P.: Dynamic partial-order reduction for modeling checking software. In: POPL. pp. 110–121. ACM Press, New York, NY, USA (2005)
9. Gray, I., Audsley, N.: Targeting complex embeeded architectures by combining the multicore communications API (mcapi) with compile-time virtualisation. In: LCTES. ACM, Chicago, Illinois, USA (2011)
10. Lahiri, S.: SMT-based modular analysis of sequential systems code. In: CAV. Springer-Verlag (2011)
11. Lahiri, S., Qadeer, S.: Back to the future revisiting precise program verification using SMT solvers. In: POPL. ACM, San Francisco, California, USA (2008)
12. Morse, E., Vrvilo, N., Mercer, E., McCarthy, J.: Modeling asynchronous message passing for C program. In: VMCAI. vol. 7148 of LNCS (2012)
13. de Moura, L., Bjørner, N.: Z3: An efficient SMT solver. In: TACAS. vol. 4963, pp. 337–340. Springer, Heidelberg (2008)
14. Netzer, R., Brennan, T., Damodaran-Kamal, S.: Debugging race conditions in message-passing programs. In: ACM SIGMETRICS Symposium on Parallel and Distributed Tools. pp. 31–40. Philadelphia, PA, USA (1996)
15. Sharma, S., Gopalakrishanan, G., Mercer, E., Holt, J.: Mcc - a runtime verification tool for mcapi user applications. In: FMCAD (2009)
16. Vakkalanka, S., Vo, A., Gopalakrishnan, G., Kirby, R.: Reduced execution semantics of mpi: From theory to pratice. In: FM. pp. 724–740 (2009)
17. Wang, C., Chaudhuri, S., Gupta, A., Yang, Y.: Symbolic pruning of concurrent program executions. In: ESEC/FSE. pp. 23–32. ACM, New York, NY, USA (2009)

18. Zhang, Y., Evelyn, D.: Barrier matching for programs with textually unaligned barriers. In: PPoPP. pp. 194–204. ACM, San Jose, California, USA (2007)