

Proving MCAPI Executions are Correct

Applying SMT Technology to Message Passing

Abstract

Asynchronous message passing for C language is an important paradigm in writing applications for embedded heterogeneous multicore systems. The Multicore Association (MCA), an industry consortium promoting multicore technology, is working to standardize message passing into a single API, MCAPI, for bare metal implementation and portability across platforms. Correctness in such an API is difficult to reason about manually, and testing against reference solutions is equally difficult as reference solutions implement an unknown set of allowed behaviors, and programmers have no way to directly control API internals to expose or reproduce errors. This paper encodes an MCAPI execution as an SMT problem, which if satisfiable, yields a feasible execution schedule on the same trace, such that it resolves non-determinism in the MCAPI runtime in the way that it now fails user provided assertions. Such an encoding is useful for test, debug, and verification of MCAPI program execution. The novelty in the encoding is the use of match pairs (potential send and receive couplings) to model the MCAPI execution in the SMT problem. Match-pair encoding for MCAPI executions, when compared to other encoding strategies, is simpler to reason about, results in significantly fewer terms in the SMT problem, and captures feasible behaviors that are ignored in previously published techniques.

Keywords Abstraction, refinement, SMT, message passing

1. Introduction

Embedded devices fill all sorts of crucial roles in our lives. They exist as medical devices, as network infrastructure, and they control our automobiles. Embedded devices continue to become more powerful as computing hardware becomes smaller and more modular. It is now becoming commonplace to find multiple processing units inside a single device. The Multicore Association (MCA) [2] is an industry group that has formed to define specifications for low-level communication, resource management, and task management for multicore devices.

One specification that the MCA has released is the Multicore Association Communications API (MCAPI) [3]. The specification defines basic data type, data structures, and functions that can be used to perform simple message passing operations between different computing entities within a device. Messages can be passed across persistent channels that force an ordering of the messages, or they can be passed to specific *endpoints* within the system. The

specification places few ordering constraints on messages passed from one endpoint to another. This freedom introduces into the system possibilities that two or more messages are racing if the order of their arrival at the destinations are non-deterministic [17]. Without a way to explore this non-determinism in the MCAPI runtime, it is not possible to test and debug the programs.

Sharma et al. created a method of using concrete execution to verify MCAPI programs, but it was later discovered that this method does not completely explore the entire execution space of certain kinds of MCAPI programs [18]. The method provides match pairs – couplings for potential sends and receives that we use in our work. Instead of exploring all relevant interleavings of a program in the concrete execution in [18], Wang et al. provided a symbolic algorithm that verifies each partitioned concurrent trace program (CTP) with shared memory semantics using a satisfiability solver [20]. Elwakil et al. also defined a method of representing MCAPI program executions as Satisfiability Modulo Theory (SMT) problems building on the method of [20] adapting it to message passing, but this method does not correctly model all kinds of MCAPI program executions [9]. A shortcoming of this method is the assumption that the user is able to provide the exact set of match pairs. Such an assumption is not reasonable for large complex program traces. In this paper, however, we assume a user provides a possible set of match pairs as input to an MCAPI program. As such, the existing technique discussed above either faces the state explosion problem [18], or provides incorrect encoding that does not capture the true program behavior in an MCAPI runtime [9]. In order to prevent those problems, this paper discusses a new method that provides both efficiency and correctness.

We present an encoding for MCAPI program executions that uses match pairs from [18] in a new SMT encoding that requires fewer terms than [9] but includes all possible program executions unlike [9]. In other words, the solution for our encoding provides efficiency. Most importantly, this method correctly captures all possible execution traces allowed by the MCAPI specification in a trace of a concurrent program as an SMT problem enabling the exploration of inherent non-determinism that may exist in any MCAPI runtime implementation. Our main contributions in this paper include:

1. A correct and efficient SMT encoding of an MCAPI program that guarantees of detecting all program errors if the user provided match pairs are precise or over-approximated.
2. An $O(N^2)$ algorithm to generate an over-approximation of possible match-pairs, where N is the code line numbers of a program.

The rest of the paper is organized as follows: Section 2 presents an MCAPI program in which two program traces exist due to non-determinism in the MCAPI runtime. Section 3 defines an SMT encoding of an MCAPI program execution that captures non-determinism in message order by defining the set of possible match pairs between sends and receives on a common endpoint. It is

proved that an over-approximated match set in our SMT encoding of an MCAPI program can reflect the actual program executions as the true set does. Section 4 solves the outstanding problem in other encodings of generating feasible match pairs to use in the encoding. It presents a $O(N^2)$ algorithm that over-approximates the precise match set, where N is the code line numbers of a program. Section 5 presents the experimental results that show our encoding to be correct and efficient, as other existing encodings both omit critical program behaviors and require more SMT clauses. Section 6 discusses related work, and Section 7 presents our conclusions and future work.

2. Example

It is a challenge to explain intended behavior in simple scenarios consisting of a handful of calls when dealing with concurrency. Consider the MCAPI program in Figure 1 that includes three tasks that use send (`mcapi_msg_send_i`) and receive (`mcapi_msg_recv_i`) calls to communicate with each other. Line numbers appear in the first column for each task, and the declarations of the local variables are omitted for space. Picking up the scenario just after the endpoints are defined, lines 02 and 05 of task 0 receive two messages on the endpoint `e0` in variables `A` and `B` which are converted to integer values and stored in variables `a` and `b` on lines 04 and 07; task 1 receives one message on endpoint `e1` in variable `C` on line 03 and then sends the message “1” on line 05 to `e0`; and finally, task 2 sends messages “4” and “Go” on lines 04 and 06 to endpoints `e0` and `e1` respectively. Task 0 has additional code (lines 08 - 09) to assert properties of the values in `a` and `b`. The `mcapi_wait` calls block until the associated send or receive buffer is able to be used. Given the scenario, we might ask the question: “What are the possible values of `a` and `b` after the scenario completes?”

The intuitive trace is presented in the first three columns of Figure 2. Note that the first column contains the line number of each command shown in the trace. The second column contains the task number in Figure 1. The third column presents the commands identified by the source line numbers shown in Figure 1. Also, we define a shorthand for each command of send (denoted as S), receive (denoted as R), or wait (denoted as W) in the fourth column for convenience in the presentation. For each command $0_{i,j}(k, \&h)$, $0 \in \{S, R\}$ or $W(\&h)$, i represents the task number, j represents the source line number, k represents the destination task number, and h represents the command handler. Note that a specific destination task number can be assigned to each receive for a provided program trace, which implies the matched send in the program runtime. To present an MCAPI program, the destination task number of each receive is replaced with “*” to provide program non-determinism. From the trace, variable `a` should contain 4 and variable `b` should contain 1 since task 2 must first send message “4” to `e0` before it can send message “Go” to `e1`; consequently, task 1 is then able to send message “1” to `e0`. The assume notation asserts the control flow taken by the program execution which in this example, asserts the true branch of the condition on line 08 of task 0. At the end of execution the assertion on line 09 of task 0 holds and no error is found. Such intuition is a valid program execution.

There are two ways to implement the MCAPI semantics for matching sends to receives including a infinite buffer semantics (the message is copied into a runtime provided buffer on the API call) or a zero-buffer semantics (the message is copied into a endpoint provided buffer) [19]. An infinite-buffer semantics provides more non-deterministic because a particular endpoint can receive more messages under the circumstance of *infinite-buffer* while the *zero-buffer* can only allow one message to be received by a specific endpoint at a time. For our example in Figure 1, if we use *zero-buffer* as the buffer setting, the send call on line 05 of task 1 cannot

match with the receive call on line 02 of task 0. This is because the send call on line 04 of task 2 must be completed before the send call on line 06 of task 2 can match with the receive call on line 03 of task 1. When using *infinite-buffer*, however, we can have another scenario shown in the fourth and fifth column of Figure 2 written in the shorthand notation. The variable `a` contains 1 instead of 4, since the message “1” is sent to `e0` after sending the message “Go” to `e1` as it is possible for the send on line 04 of task 2 to be delayed in transit. The specification indicates that the wait on line 05 of task 2 returns once the buffer is available. That only means the message is somewhere in the MCAPI runtime; it does not mean the message is delivered. As such, it is possible for the message to be delayed in transit allowing the send from task 1 on line 05 to arrive at `e0` first and be received in variable “`a`”. Such a scenario is a program execution that results in an assertion failure at line 09 in Figure 1.

From the discussion above, it is important to consider non-determinism in the MCAPI runtime when testing or debugging an MCAPI program execution. Note that we use *infinite-buffer* for the program execution in the rest of our discussion. The next section presents an encoding algorithm that takes an MCAPI program execution with a set of possible send-receive match pairs and generates an SMT problem that if satisfied proves that non-determinism can be resolved in a way that violates a user provided assertion and if unsatisfiable proves the trace correct (meaning the user assertions hold on the given execution). The encoding can be solved by an SMT solver such as Yices [7] and Z3 [16], and the non-deterministic behavior of the program can be resolved.

3. SMT Model

The novelty of the SMT encoding in this paper is its use of match pairs rather than the state-based or order-based encoding of prior work [8, 9]. The algorithm to create the encoding takes as input a set of possible match pairs and a trace through an MCAPI program with the appropriate assumes and asserts as shown in Figure 2. A match pair is the coupling of a receive to a particular send. Figure 3(a) is the set of possible match pairs for the program in Figure 1 using our shorthand notation defined in Figure 2. The set admits, for example, that $R_{0,2}$ can be matched with either $S_{1,5}$ or $S_{2,4}$. The SMT encoding in this paper asks the SMT solver to resolve the match pairs for the system in such a way that the final values of program variables meet the assumption on control flow but violate some assertion.

Before presenting our SMT encoding for a specific example, we need to explain auxiliary data structures and support functions in detail. We give the definitions in the Yices [7] input language.

```
(define HB :: (→ a::int b::int
  (subtype (c::bool) (ite (< a b) (= c true) (= c false)))))
(1)
```

The HB function in equation (1) creates a happens-before relationship between two events. The notation “ite” refers to a condition statement such that if the condition is satisfied, then the first subclause is true; otherwise, the second subclause is true. The function takes two events `a` and `b` as parameters, such that $a, b \in \mathbb{N}$, where \mathbb{N} is the set of natural numbers. In our SMT encoding, every location is assigned an event, and HB orders those locations (i.e., program order, etc.). This function creates a constraint such that the first event must be less than the second event, indicating that the first event occurs before the second event. For example, action `a` happens before `b` if and only if $event_a$ for `a` is greater than $event_b$ for `b`. This function is not only used to assert the program order of statements within the same thread, but also to assert, for any matched pair, that the send operation occurs before the receive.

Task 0	Task 1	Task 2
00 initialize(NODE_0,&v,&s);	00 initialize(NODE_1,&v,&s);	00 initialize(NODE_2,&v,&s);
01 e0 = create_endpoint(PORT_0,&s);	01 e1 = create_endpoint(PORT_1,&s);	01 t2 = create_endpoint(PORT_2,&s);
	02 e0 = get_endpoint(NODE_0,PORT_0,&s);	02 t0 = get_endpoint(NODE_0,PORT_0,&s);
		03 t1 = get_endpoint(NODE_1,PORT_1,&s);
02 msg_rcv_i(e0,A,sizeof(A),&h1,&s);	03 msg_rcv_i(e1,C,sizeof(C),&h3,&s);	04 msg_send_i(e2,e0,"4",2,N,&h5,&s);
03 wait(&h1,&size,&s,MCAPI_INF);	04 wait(&h3,&size,&s,MCAPI_INF);	05 wait(&h5,&size,&s,MCAPI_INF);
04 a = atoi(A);		
	05 msg_send_i(e1,e0,"1",2,N,&h4,&s);	06 msg_send_i(e2,e1,"Go",3,N,&h6,&s);
05 msg_rcv_i(e0,B,sizeof(B),&h2,&s);	06 wait(&h4,&size,&s,MCAPI_INF);	07 wait(&h6,&size,&s,MCAPI_INF);
06 wait(&h2,&size,&s,MCAPI_INF);		
07 b = atoi(B);	07 finalize(&s);	08 finalize(&s);
08 if(b > 0);		
09 assert(a == 4);		
10 finalize(&s);		

Figure 1. An MCAPI concurrent program

Trace 1				Trace 2			
Line	Task	Command	Shorthand	Task	Command		
0	2	04 msg_send_i(e2,e0,"4",2,N,&h5,&s);	S _{2,4} (0,&h5)	2	04 S _{2,4} (0,&h5)		
1	2	05 wait(&h5,&size,&s,MCAPI_INF);	W(&h5)	2	05 W(&h5)		
2	0	02 msg_rcv_i(e0,A,sizeof(A),&h1,&s);	R _{0,2} (2,&h1)	2	06 S _{2,6} (1,&h6)		
3	0	03 wait(&h1,&size,&s,MCAPI_INF);	W(&h1)	2	07 W(&h6)		
4	2	06 msg_send_i(e2,e1,"Go",3,N,&h6,&s);	S _{2,6} (1,&h6)	1	03 R _{1,3} (2,&h3)		
5	2	07 wait(&h6,&size,&s,MCAPI_INF);	W(&h6)	1	04 W(&h3)		
6	0	04 a = atoi(A);		1	05 S _{1,5} (0,&h4)		
7	1	03 msg_rcv_i(e1,C,sizeof(C),&h3,&s);	R _{1,3} (2,&h3)	1	06 W(&h4)		
8	1	04 wait(&h3,&size,&s,MCAPI_INF);	W(&h3)	0	02 R _{0,2} (1,&h1)		
9	1	05 msg_send_i(e1,e0,"1",2,N,&h4,&s);	S _{1,5} (0,&h4)	0	03 W(&h1)		
10	1	06 wait(&h4,&size,&s,MCAPI_INF);	W(&h4)	0	04 a = atoi(A);		
11	0	05 msg_rcv_i(e0,B,sizeof(B),&h2,&s);	R _{0,5} (1,&h2)	0	05 R _{0,5} (2,&h2)		
12	0	06 wait(&h2,&size,&s,MCAPI_INF);	W(&h2)	0	06 W(&h2)		
13	0	07 b = atoi(B);		0	07 b = atoi(B);		
14	0	08 assume(b > 0);		0	08 assume(b > 0);		
15	0	09 assert(a == 4);		0	09 assert(a == 4);		

Figure 2. Two execution traces of the MCAPI program in Figure 1

(define-type *R* (record *M::int e::int value::int event::int*)
(define-type *S* (record *O::int ID::int e::int value::int event::int*)
(2)

We define two records in equation (2) that are essential to encoding our SMT problem. The first record *R* defined above represents a receive action such that *M*, *e*, *value*, *event* ∈ ℕ. *M* represents a free variable (notation where substitution may take place) that is assigned the *event* field of the matched record *S*. *e* is the endpoint of the receive action *R*, *value* represents a value received by the receive operation, and *event* is the event term which indicates the program order of the receive action *R*. The second record *S* defined above is a send action such that *O*, *ID*, *e*, *value*, *event* ∈ ℕ. The fields *e* and *event*, as those in tuple *R*, are the destination endpoint and event term, respectively. The field *value* is the value being sent. *O* is a free variable used for match pairs, and it is assigned an event term of a receive operation if the receive operation is to be matched to the send operation. *ID* is a unique number that identifies each send operation. Note that the integer value assigned to *event* indicates the program order. Also, we use two match variables, *M* and *O*, for different purposes. *M* is used for preventing two receive operations from being matched with the same send operation. *O* is used for ordering two receive operations with respect to the program order of two matched send operations. The SMT solver must assign the *M* and *event* fields of the send tuple, and it must assign the *event*, *M*, and *value* fields of the receive tuple in a way that is consistent with all the constraints in the encoding.

(define MATCH :: (→ *r::R s::S* (subtype (*c::bool*) (ite
(and (= (select *r e*) (select *s e*))
(= (select *r value*) (select *s value*))
(HB (select *s event*) (select *r event*))
(= (select *r M*) (select *s ID*))
(= (select *s O*) (select *r event*))))
(= *c* true)
(= *c* false))))

The MATCH function in equation (3) takes tuples *R* and *S* as parameters. Each “select” statement selects a field of tuple *R* or *S*. The function creates a number of constraints that represent the pairing of the specified send and receive operations. This function asserts that the destination endpoint of the send operation is the same as the endpoint used by the receive operation, that the value sent by the send operation is the same value received by the receive operation, that the send operation occurs before the receive operation, that the “match” value (*M*) in the receive tuple is equal to the “ID” value in the send tuple, and that the “O” value in the send tuple is equal to the “event” value in the receive tuple. Note that the encoding is going to use the “O” value in conjunction with the HB function to enforce message non-overtaking when messages are sent from the same endpoint. Consider a simple example below

that sends two messages from a Task 0 to Task 1,

Task 0	Task 1
$S_{0,1}(1, \&h1)$	$R_{1,1}(*, \&h3)$
$S_{0,2}(1, \&h2)$	$R_{1,2}(*, \&h4)$
$W(\&h1)$	$W(\&h3)$
$W(\&h2)$	$W(\&h4)$

The “O” values in the send records will be assigned to the order tracking events for $R_{1,1}$ and $R_{1,2}$ in the final encoding. Message non-overtaking from transactions on common endpoints is encoded by forcing the send events to be received in program order using the HB function as below in our simple example:

(HB (select $S_{0,1}$ O) (select $S_{0,2}$ O))

(define NE :: ($\rightarrow r1::R\ r2::R$ (subtype ($c::bool$) (ite
 (= (select $r1\ M$) (select $r2\ M$))
 (= c false)
 (= c true)))))) (4)

The NE function in equation (4) is used to assert that no two receive operations are matched with the same send operation. The parameters for this function are two receive tuples. An assertion is made that the values of their “match” fields, M , are not equal. This shows that they are paired with two different send operations.

Prior SMT models of MCAPI program executions implicitly compute match pairs by adding possible happens-before relations on sends and receives with the conditions under which those orderings are valid. The SMT solver is then asked to resolve the happens-before relation by choosing specific orders of sends and receives. Match pair encoding, as presented in this work, though it has the same computational complexity as order based encoding (you still need to figure out match pairs on endpoints), is simpler to reason about directly rather than implicitly through orders, results in significantly fewer terms in the SMT problem, and does not restrict out possible matches allowed by the MCAPI specification of the runtime as do existing encodings.

The presentation of the SMT encoding is structured as

$smt = (defs\ constraints\ match)$

where the field *defs* represents all definitions of the send, receive operations and the free variables; The field *constraints* represents all constraint clauses, including the assert clauses and the clauses built by the HB functions; and the field *match* represents the clauses built by the MATCH function on a set of match pairs and the NE function.

The *Intra-Happens-Before-Order* (Intra-HB) in [19] constrains the Happens-Before relations between commands in any control flow path. Those relations are essential to their algorithm for their Message-passing verifier. Our SMT structure encodes all the relations in *Intra-HB*, such that it can follow any control flow path executed by the algorithm in [19]. The *Intra-HB* consists of 8 *Happens-Before* relations:

$\prec_{mhb} = \prec_S \cup \prec_R \cup \prec_W \cup \prec_{SW} \cup \prec_{RW} \cup \prec_{mb} \cup \prec_{fo} \cup \prec_{io}$

The relations for *Barrier* commands are omitted because they are not supported by MCAPI. The relations \prec_S , \prec_R and \prec_W are used to constrain the state transition of commands. For example, issuing a send command happens before matching it with some receive in transit. Instead of explicitly giving the state information, the SMT structure in this paper implicitly encodes the state transitions in the *constraints* area of the encoding. The relations \prec_{SW} , \prec_{RW} , \prec_{fo} and \prec_{io} constrain the program order of commands. We encode those relations in the field *constraints* of our SMT encoding. The relation \prec_{mb} builds the match before for sends and receives. Also, we have the relation in the field *match* in our SMT encoding.

defs is not shown;
constraints;
 00 (HB $R_{0,2}$.event $W(\&h1)$.event)
 01 (HB $W(\&h1)$.event $R_{0,5}$.event)
 02 (HB $R_{0,5}$.event $W(\&h2)$.event)
 03 (HB $W(\&h2)$.event $assume$.event)
 04 (HB $assume$.event $assert$.event)
 05 (HB $R_{1,3}$.event $W(\&h3)$.event)
 06 (HB $W(\&h3)$.event $S_{1,5}$.event)
 07 (HB $S_{1,5}$.event $W(\&h4)$.event)
 08 (HB $S_{2,4}$.event $W(\&h5)$.event)
 09 (HB $W(\&h5)$.event $S_{2,7}$.event)
 10 (HB $S_{2,7}$.event $W(\&h6)$.event)
 11 (assert ($> b\ 0$))
 12 (assert (not ($= a\ 4$)))
match;
 13 (or (MATCH $R_{0,2}\ S_{2,4}$)
 14 (MATCH $R_{0,2}\ S_{1,5}$)
 15 (or (MATCH $R_{0,5}\ S_{2,4}$)
 16 (MATCH $R_{0,5}\ S_{1,5}$)
 17 (MATCH $R_{1,3}\ S_{2,7}$)
 18 (NE $R_{0,2}\ R_{0,5}$))

(a)

(b)

Figure 3. A match pair set and SMT encoding of the system in Figure 1. (a) The match pairs based on endpoints. (b) The SMT encoding where HB creates a happens-before constraint, MATCH creates a match pair constraint, and NE creates a constraint to preclude the use of conflicting match pairs.

The input for generating the SMT encoding is an MCAPI program together with a set of match pairs. The first trace in Figure 2 with the set of match pairs in Figure 3(a) is such an example. Given an input, the algorithm for generating our encoding is enumerated as follows. For convenience, we use the notation “*.event*” following each command to represent the event field in our encoding.

1. Define variable “*.event*” for the program order of each assume or assert command in the field *defs*;
2. for each assume, add a statement to *constraints*, and for each assert, add a negated assert to *constraints* in order to obtain an execution trace with assertion failure from an SMT solver;
3. for each send or receive in the trace, create an *S* or *R* definition and create an associated “*.event*” variable for program order and add these to *defs*;
4. for each program order represented by the field *event* in each thread, add an HB relation on the associated event for program order to *constraints*;
5. for each match pair for a send and receive as input, add an MATCH relation to *match*;
6. for two sends $S_{i,j}$ and $S_{i,k}$ in the same task i sent to an identical endpoint such that $i < k$, if the relation (HB $S_{i,j}$.event $S_{i,k}$.event) is defined in *constraints*, add the clause (HB (select $S_{i,j}\ O$) (select $S_{i,k}\ O$)) to *constraints*.
7. collect into a disjunction for all possible send matches for any given receive in the field *match*;
8. and collect into a NE relation for all receives that are matched with any identical send and add it to *match*.

We use our execution trace in Figure 2 (left side) with the set of match pairs in Figure 3(a) as input to go through the algorithm above for generating an SMT encoding shown in Figure 3(b). Note that we use the wait command with handler in Figure 2 to denote the wait command for a send or receive command. First, we use steps 1 and 3 of the algorithm above to generate the *defs* “area” that is not shown because the definitions are not novel to our solution.

Second, we use step 2 to generate the asserts shown on line 11 and 12 for the assume and assert commands of the original execution trace in Figure 2. The first assert on line 11 works as the assume (line 14 at trace 1) in Figure 2, such that it prevents the SMT solver from finding solutions that are not consistent with control flow which requires “ $b \geq 0$ ”. The second assert on line 12 is negated as the goal is to find schedules that violate the property. Third, we use step 4 to generate line 00 - 10, which represent the HB relationship between the events of commands. By encoding function HB, $R_{0,2}$ happens before the wait command $W(\&h1)$ given the “event” for each command as parameters on line 00 in Figure 3(b). Note that any assignment to event variables by the SMT solver must comply with program order constraints. Each event in the HB orderings on line 00 - 10 are ordered in a thread of Figure 2. Finally, we use steps 5, 7 and 8 to generate the *match* “area” of the SMT encoding. In particular, we send the set of match pairs in Figure 3(a) to the algorithm and use step 5 to generate an MATCH statement for each match pair; collect into a disjunction by step 7 so that all possible MATCH statements are organized on line 13 - 17; add NE relation by step 8 that precludes $R_{0,2}$ and $R_{0,5}$ from matching to the same send on line 18.

Other than the basic structure of the SMT encoding, we use the function

$$\text{ANS}(smt) \mapsto \{\text{SAT}, \text{UNSAT}\}$$

to return the solution of an SMT problem, such that **SAT** represents a satisfiable solution that finds an execution of the MCAPI program that violates the user defined correctness property, and **UNSAT** represents an unsatisfiable solution that all possible execution traces either meet the correctness property in the same control flow, or follow a different control flow. Note that **UNSAT** and **SAT** are ordered such that **UNSAT** < **SAT**.

The SMT encoding defined above is used to capture the non-deterministic behavior of an MCAPI program by giving a complete set of match pairs. As we discussed in the previous section, the MCAPI program in Figure 1 contains two outcomes of execution as defined in the MCAPI specification. The SMT encoding we present in Figure 3(b) captures both executions, since the set of match pairs in Figure 3(a) is a complete set where all matches that can occur in the real execution for our running example in Figure 1 are included, and all matches that cannot occur in the real execution are not included. Furthermore, the following theorem states that we can over-approximate the true set of match pairs and still prove correctness. If there is no error with the over-approximated set, then there is no error arising from non-determinism in the runtime on that program execution. If there is an error from the over-approximated set, that error is also guaranteed a real error in the program runtime by giving the second theorem in the following. Note that two SMT problems smt_α and smt_β in the following theorem have an identical *defs* and *constraints* sets, and the match set of smt_α is the subset of that of smt_β so that smt_β represents an over-approximation of smt_α .

3.1 Theorem 1

For two SMT problems, $smt_\alpha = (\text{defs constraints match}_\alpha)$ and $smt_\beta = (\text{defs constraints match}_\beta)$, if $\text{set}(\text{match}_\alpha) \subseteq \text{set}(\text{match}_\beta)$,

$$\text{ANS}(smt_\alpha) \leq \text{ANS}(smt_\beta)$$

where $\text{set}(\text{match})$ represents the input set of match pairs for the MATCH clauses in the field *match*.

Proof Sketch. Consider the MCAPI program in Figure 1 as an example. Figure 4(a) and (b) are two different SMT encodings for our running example in Figure 1 generated from different sets of possible match pairs, such that Figure 4(a) encodes trace 1 in Figure 2 and Figure 4(b) encodes trace 2 in Figure 2. By solving the

<i>defs</i> is not shown;	<i>defs</i> is not shown;
<i>constraints</i> ;	<i>constraints</i> ;
00 (HB $R_{0,2}$.event $W(\&h1)$.event)	00 (HB $R_{0,2}$.event $W(\&h1)$.event)
01 (HB $W(\&h1)$.event $R_{0,5}$.event)	01 (HB $W(\&h1)$.event $R_{0,5}$.event)
02 (HB $R_{0,5}$.event $W(\&h2)$.event)	02 (HB $R_{0,5}$.event $W(\&h2)$.event)
03 (HB $W(\&h2)$.event assume.event)	03 (HB $W(\&h2)$.event assume.event)
04 (HB assume.event assert.event)	04 (HB assume.event assert.event)
05 (HB $R_{1,3}$.event $W(\&h3)$.event)	05 (HB $R_{1,3}$.event $W(\&h3)$.event)
06 (HB $W(\&h3)$.event $S_{1,5}$.event)	06 (HB $W(\&h3)$.event $S_{1,5}$.event)
07 (HB $S_{1,5}$.event $W(\&h4)$.event)	07 (HB $S_{1,5}$.event $W(\&h4)$.event)
08 (HB $S_{2,4}$.event $W(\&h5)$.event)	08 (HB $S_{2,4}$.event $W(\&h5)$.event)
09 (HB $W(\&h5)$.event $S_{2,7}$.event)	09 (HB $W(\&h5)$.event $S_{2,7}$.event)
10 (HB $S_{2,7}$.event $W(\&h6)$.event)	10 (HB $S_{2,7}$.event $W(\&h6)$.event)
11 (assert (> b 0))	11 (assert (> b 0))
12 (assert (not (= a 4)))	12 (assert (not (= a 4)))
<i>match</i> ;	<i>match</i> ;
13 (MATCH $R_{0,2}$ $S_{2,4}$)	13 (MATCH $R_{0,2}$ $S_{1,5}$)
14 (MATCH $R_{0,5}$ $S_{1,5}$)	14 (MATCH $R_{0,5}$ $S_{2,4}$)
15 (MATCH $R_{1,3}$ $S_{2,7}$)	15 (MATCH $R_{1,3}$ $S_{2,7}$)
(a)	(b)

Figure 4. Two SMT encodings of the system in Figure 1. (a) The SMT encoding for Trace 1 in Figure 2. (b) The SMT encoding for Trace 2 in Figure 2.

encodings in Figure 4(a) and (b) for trace 1 and 2 in Figure 2 respectively, we get an unsatisfiable solution for Figure 4(a), and a satisfiable solution for Figure 4(b). As we discussed in Section 2, trace 1 is an execution without failure of the assertion, and trace 2 is the one that fails the assertion. Compare both encodings with that in Figure 3(b), we find that the fields *defs* and *constraints* are identical except for the MATCH clauses. In particular, the input set of match pairs of either Figure 4(a) or (b) is the subset of that in Figure 3(b). As discussed above, the encoding in Figure 3(b) captures the non-deterministic behavior of our running program in Figure 1, which encodes trace 1 and 2 in Figure 2 into one single SMT problem. Thus, an SMT solver will return a satisfiable solution for the encoding in Figure 3(b). Thus, ANS on the encoding in Figure 3(b) is greater than or equal to the ANS on the encoding in either Figure 4(a) or (b). The formal proof is given in a long version of paper at (“<http://students.cs.byu.edu/yhuang2/downloads/paper.pdf>”).

Given $\text{set}(\text{match}_\alpha)$ and $\text{set}(\text{match}_\beta)$ in the theorem above a complete set of match pairs and an over-approximated set for an MCAPI program, respectively, we can further prove that $\text{ANS}(smt_\alpha) = \text{ANS}(smt_\beta)$ by giving the following theorem. Note that a match pair $(R, S) \in \text{set}(\text{match}_\beta) / \text{set}(\text{match}_\alpha)$ is called “bogus”, since it cannot exist in a real execution of the program.

3.2 Theorem 2

Any match pair (R, S) used in a satisfying assignment of an SMT encoding smt is a valid match pair and reflects an actual possible MCAPI program execution.

Proof. Proof by contradiction. Suppose (R, S) is a “bogus” match pair that causes $\text{ANS}(smt) = \text{SAT}$. Since (R, S) is not a valid match pair, it cannot reflect an actual program execution. In other words, the HB constraints encoded in smt are not satisfied. Based on the fact above, the answer of smt is **UNSAT** and it contradicts the previous hypothesis. Thus, (R, S) is a valid match pair in smt and reflects an actual possible MCAPI program execution. \square

By proving Theorem 2, we infer that a “bogus” match pair can only cause an unsatisfying assignment of an SMT problem. Further, given that $\text{set}(\text{match}_\alpha)$ and $\text{set}(\text{match}_\beta)$ reflect a complete set and an over-approximated set respectively, the answers of two encodings smt_α and smt_β discussed above are equal since any “bogus” match pair involved in smt_β is only used in unsatisfying assignments.

```

//initialization
input an MCAPI program
initialize list_r by traversing each task of the program and storing
each receive command
initialize list_s by traversing each task of the program and storing
each send command

//check each receive and send with the same endpoint
for r in list_r
  for s in list_s
    //check matching criteria for r and send
    if
      1.the endpoint of r is equal to the destination endpoint of s, and
      2.the index of r is larger or equal to the index of s, and
      3.the index of r is less or equal to the index of s
      plus the size of list_s minus the number of sends with the same
      source and destination endpoints of s
    then
      add pair (r, s) to match_set
    else
      do next iteration
    end if
  end for
end for

output match_set;

```

Figure 5. Pseudocode for generating over-approximated match pairs

The non-determinism of an MCAPI program discussed above relies on the *infinite-buffer* setting. If the *zero-buffer* setting is used instead, the SMT encoding needs to be reorganized as follows. For each match pair that cannot exist in the program runtime under the *zero-buffer* setting, add extra HB clauses in the SMT encoding to prevent incorrect behavior. For example in Figure 1, to prevent $R_{1,3}$ to be matched with $S_{2,6}$ before $R_{0,2}$ is matched with $S_{2,4}$ prohibited by the *zero-buffer*, an HB relation is added such that the wait command $W(\&h1)$ for $R_{0,2}$ happens before $S_{0,6}$.

4. Generating Match Pairs

The exact set of match pairs can be generated by simulating the program trace and using a depth-first search to enumerate non-determinism arising from concurrent sends and receives. Such an effort, however, solves the very problem we are trying to solve because if you simulate the program trace exploring all non-determinism, then you may as well verify all runtime choices for property violations at the same time. In this section, we present an algorithm that does not require an exhaustive enumeration of runtime behavior in simulation. Our algorithm over-approximates the match pairs such that match pairs that can exist in the runtime are all included and some bogus match pairs that cannot exist in the runtime may or may not be included.

The algorithm generates the over-approximated match pair set by matching each pair of the send and receive commands at common endpoints and then pruning obvious matches that cannot exist in any runtime implementation of the specification. Figure 5 presents the major steps of the algorithm. Algorithm proceeds by first linearly traversing each task of the program storing each receive and send command into two distinct structured lists. Send list $list_s$ is structured as in (6) and receive list $list_r$ is structured as in (5).

$$\begin{aligned}
(e_0 &\rightarrow ((0, R_{0,1}), (1, R_{0,2}), \dots)) \\
(e_1 &\rightarrow ((0, R_{1,1}), (1, R_{1,2}), \dots)) \\
&\dots \\
(e_n &\rightarrow ((0, R_{n,1}), (1, R_{n,2}), \dots))
\end{aligned} \tag{5}$$

The list in (5) keeps record of the receive commands uniquely identified by the endpoint e_x , where x is the number of the endpoint. The integer in the first field of each pair indicates the program order of the receive commands within each task on the specified endpoint. Note that the receive command with lower program order should be served first in the program runtime. The receive command in the

Task 0	Task 1	Task 2
$R_{0,1}(*, \&h1)$	$S_{1,1}(0, \&h5)$	$S_{2,1}(0, \&h8)$
$R_{0,2}(*, \&h2)$	$R_{1,2}(*, \&h6)$	
$S_{0,3}(1, \&h3)$	$S_{1,3}(0, \&h7)$	
$R_{0,4}(*, \&h4)$		

Figure 6. Another MCAPI concurrent program

second field of the list is defined as a record R in equation 2 of the previous section. The notations $R_{0,1}$, $R_{0,2}$, etc. represent the unique identifiers for the receives. Recall that we operate on a program trace so there are no loops and all instances of send and receive operations are uniquely identified.

$$\begin{aligned}
&\text{“dst”} \quad \text{“src”} \quad \text{“src”} \\
(e_0 &\rightarrow ((e_1 \rightarrow ((0, S_{1,1}), (1, S_{1,2}), \dots), (e_2 \rightarrow (\dots, \dots)))) \\
(e_1 &\rightarrow ((e_0 \rightarrow ((0, S_{0,1}), (1, S_{0,2}), \dots), (e_2 \rightarrow (\dots, \dots)))) \\
&\dots \\
(e_n &\rightarrow ((e_0 \rightarrow ((0, S_{0,3}), (1, S_{0,4}), \dots), (e_1 \rightarrow (\dots, \dots))))
\end{aligned} \tag{6}$$

Also, we keep record of the send commands using the second list in (6) that is uniquely identified by the destination endpoint e_x . Each sublist for the destination endpoint is uniquely identified by the source endpoint e_y . Similarly, the integer in the first field of each pair indicates the program order of the send command within the same task from a common source and to a common endpoint. The send command in the second field of the list is defined as a record S in equation 2 of the previous section. The notations $S_{1,1}$, $S_{1,2}$, etc. represent the unique identifiers for the sends.

Figure 6 is an example program in our shorthand notation to present our algorithm. The sends $S_{1,1}$, $S_{1,3}$ and $S_{2,1}$ have Task 0 as an identical destination endpoint. The send $S_{0,3}$ has Task 1 as the destination endpoint. In our running example in Figure 6, we generate our receive list and send list in equation (7) and (8), respectively.

$$\begin{aligned}
(0 &\rightarrow ((0, R_{0,1}), (1, R_{0,2}), (2, R_{0,4}))) \\
(1 &\rightarrow ((0, R_{1,2})))
\end{aligned} \tag{7}$$

$$\begin{aligned}
(0 &\rightarrow ((1 \rightarrow ((0, S_{1,1}), (1, S_{1,3})), (2 \rightarrow ((0, S_{2,1})))))) \\
(1 &\rightarrow ((0 \rightarrow ((0, S_{0,3}))))))
\end{aligned} \tag{8}$$

Note that $R_{0,1}$ is the first receive operation in endpoint 0, and $S_{2,1}$ is the first send from the source endpoint 2 to the destination endpoint 0.

The second step for our algorithm is to linearly traverse both lists to generate match pairs between send and receive commands. Observe that the sends defined with an identical destination endpoint and an identical source endpoint are processed in a FIFO order according to the MCAPI runtime specification. The same is true of receives on a common endpoint in a common task. The FIFO ordering, sometimes referred to as message non-overtaking, lets us prune match pairs that obviously cannot be generated in any valid implementation of the MCAPI runtime.

Let I_r and I_s be the program order indication in the lists of (5) and (6) respectively and similarly R and S be the corresponding actions. Note that I_r and I_s are increased by 1 for two consecutive receives or sends, respectively. Further, N_s is the number of send actions in the program and $n_s(src, dst)$ be the number of sends from src to dst , where src and dst represent the source and destination endpoints, respectively. A send and received can be matched if and only if

1. the destination endpoint of S is identical to the endpoint of R ;
2. $I_r \geq I_s$;
3. and $I_r \leq I_s + (N_s - n_s(src, dst))$.

If all rules of the criteria are satisfied for a send and a receive, we build a match pair for the send and the receive in the result set `match_set`. In our concrete example, $R_{0,1}$ is matched with $S_{1,1}$ or $S_{2,1}$, but it cannot be matched with $S_{1,3}$ since the second rule is not satisfied such that the order for $R_{0,1}$ is less than the order of $S_{1,3}$. We repeatedly apply the criteria to match sends and receives until we completely traverse the lists in (5) and (6). The generated set of match pairs for our example in Figure 6 is imprecise such that some extra match pair, we call “bogus”, that cannot exist in the real execution of program are included. In particular, the match pair $(S_{2,1} R_{0,4})$ is a “bogus” match pair because it is not possible to order $S_{1,3}$ before $R_{0,2}$ since $R_{1,2}$ can only match with $S_{0,3}$ that must occur after $R_{0,2}$. Fortunately, the encoding in this paper is strong enough to preclude that bad match pair in any satisfying solution.

Now let us analyze the time complexity of the algorithm. Traversing the tasks linearly takes $O(N)$ to complete, where N is the total lines of code of the program. Traversing the list of receives and the list of sends takes $O(mn)$ to complete, where m is the total number of sends and n is the total number of receives. Note that $m + n \leq N$. Totally, the algorithm takes $O(N + mn) \leq O(N + N^2) \approx O(N^2)$ to complete.

As discussed in the previous section, an over-approximated match set can give the same solution of our SMT encoding for an MCAPI program as a complete set does. As our match pair generation algorithm over-approximates the true set of match pairs allowed by the runtime for an MCAPI program trace, when used with our SMT encoding it cannot result in false positives (i.e., errors arising from match pairs that cannot happen in any implementation of the MCAPI specification).

5. Experiments and Results

We have implemented the tracking and analysis algorithms in the setting of MCAPI reference implementation. Our tool is capable of handling C programs for MCAPI semantics using the Linux *PThreads* library. In particular, the tool generates a real execution trace under the environment of MCAPI reference implementation at the beginning. The match set generation algorithm presented in Section 4 then builds an over-approximated match set. The tool then builds an SMT encoding based on the generated execution trace and match set. We use the *Yices* SMT solver [7] to check the satisfiability of the generated SMT encoding.

To evaluate the reliability and efficiency of our system, we compare to [9] even though it misses valid MCAPI runtime behavior because it is the only other encoding for MCAPI programs and it suffices to illustrate the efficiency of our encoding. Two sets of benchmarks are used in the comparison. The first set consists of four “toy” examples, including the example in Figure 1. Those examples contain small amount of operations and simple message communications. Because the runtime differences for both encodings are very small, we use the number of clause from each encoding for the same example to compare the efficiency. The other set of benchmarks consists of five large programs. One of them, we call “leader election”, basically elects a leader from several candidates by message communications. The other four programs, are all extensions to the original program in Figure 1. In particular, extra iterations are added to the original program so that more message passing will occupy and more properties will be checked. In all examples, the correctness properties are numerical assertions over variables. The experiments were conducted on a PC with 1.6 GHz Intel Quad Core processor and 8GB memory running Ubuntu 14.

As for the first set of benchmarks, we encodes 17 clauses for modeling the program in Figure 1, omitting the definition of variables at the beginning of the encoding. Based on the encoding rules in [9], however, there are 47 clauses for the same program, omitting the definitions as well. Table 1 shows the number of clauses that

	Program Order	Matches	Assume&Assert	Extra Clauses
Our Encoding	11	4	2	0
Encoding in [9]	22	13	3	8

Table 1. Comparison of two encodings for the MCAPI program in Figure 1.

	Example in Figure 1	small1	small2	small3
Our Encoding	17	8	4	11
Encoding in [9]	46	33	18	44

Table 2. Comparison of two encodings for four “toy” MCAPI programs with clause number.

Test Programs		Our Encoding		Encoding in [9]	
Name	# Messages	Time(ms)	Memory	Time(ms)	Memory
LeaderElection	30				
5iterations	15				
6iterations	18				
7iterations	21				
8iterations	24				

Table 3. Comparison of two encodings for five “large” MCAPI programs with runtime and used memory.

constrain program order, match pairs, assume/asserts, and any other extra uncategorized clauses. Program order comprises the bulk of the encoding. The extra clauses in [9] come from auxiliary variables in the encoding.

Other than the program in Figure 1, we manually generate three programs and provide the experimental result in Table 3. Small1 is a program with two tasks, where two sends are in the first task, and two receives are in the second task. Small2 is a program with only one task, where one send and one receive are in this task. Small3 is a program with three tasks, where three sends are in the first task, two receives are in the second task, and one receive is in the third task. The number of clauses follows the performance of the SMT solver with fewer clauses leading to better runtime. Table 3 shows that our encoding generates 1/5 – 1/3 the clauses of that of the encoding in [9]. We believe these reductions will generalize to any given program trace and match pair set.

As for the other set of benchmarks, we compare the runtime and consumed memory for each example.

6. Related Work

Morse et al. provided a formal modeling paradigm that is callable from C language for the MCAPI interface [15]. This model correctly captures the behavior of the interface and can be applied for model checking analysis for C programs that use the API.

In [19], Vakkalanka et al. provided POE, a DPOR algorithm [10] of their dynamic verifier, ISP, for MPI programs. POE explores all interleavings of an MPI program, and is able to detect hidden deadlocks under *zero-buffer* setting. If *infinite-buffer* is applied, however, some interleavings may not be found. In [18], Sharma et al. provided MCC, a dynamic verifier for MCAPI programs. MCC systematically explores all interleavings of an MCAPI program by concretely executing the program repeatedly. MCC uses DPOR [10] to reduce the redundant interleavings of the execution, however, it does not include all possible traces allowed by the MCAPI specification. Wang et al. provided a symbolic algorithm that detects errors in all feasible permutations of statements in an execution trace in the shared memory system [20]. In this method, the program is partitioned into several concurrent trace programs

(CTPs), and the encoding for each CTP is verified using a satisfiability solver. Elwakil et al. provided a similar work with ours [8, 9]. In their work, the method of [20] is used and adapted to the message passing system. As shown in the previous section of this paper, their method does not correctly encode all possible execution traces of an MCAPI program.

The SMT/SAT based Bounded Model Checking is one avenue of verifying and debugging systems. Burckhardt et al. presented CheckFence prototype in [5], which exhaustively checks all executions of a test program by translating the program implementation into SAT formulas. The standard SAT solver can construct counterexample if incorrect executions exist. Instead of over-approximating at the beginning and then further compressing the observations, CheckFence in [5] increments the observations each time step by adding constraints to SAT formulas. Dubrovin et al. provides a method in [6] that translates an asynchronous system into a transition formula for three partial order semantics. Other than adding constraints to the SAT/SMT formulas in order to compress the search space, the method in [6] decreases the search bound by allowing several actions to be executed simultaneously within one step. Kahlon et al. presented a Partial Order Reduction method called *MPOR* in [12]. This method cannot only be applied to the explicit state space search as other partial order reduction methods do, but also can be applied to the SMT/SAT based Bounded Model Checking. *MPOR* guarantees that exactly one execution is calculated per each Mazurkiewicz trace, in order to reduce the search space. There are several applications of the SMT/SAT based Bounded Model Checking. [14] presented a precise verification of heap-manipulating programs using SMT solvers. [13] presented some challenges in SMT-based verification for sequential system code, and tackled these issues by extending the standard SMT solvers.

The application of static analysis is another interesting thread of research to test or debug message passing programs. [21] and [4] are the approaches for MPI [1], another message passing interface standard. [11] presented a system that uses static analysis to determine offline the topology of the communications and nodes in the input MCAPI program.

7. Conclusions and Future Work

We have presented an SMT encoding of an MCAPI program that uses match pairs rather than the state-based or order-based encoding in the prior work. Unlike the existing method of SMT encoding, our encoding can correctly capture the non-deterministic behavior of an MCAPI program. Also, we have provided an algorithm with $O(N^2)$ time complexity that over-approximates the true set of match pairs, where N is the total number of code lines of the program. Once the match pairs are generated by the algorithm, our encoding guarantees the actual possible MCAPI program executions can be captured. The experimental results shows that our encoding can be solved efficiently compared with the existing work in [9].

The precision of the set of match pairs is essential to the efficiency of SMT encodings. Currently we have an over-approximated generation method which still keeps several “bogus” match pairs in the generated set. New methods for generating a much more precise set of match pair are required. Also The method of DPOR are considered to combine with our SMT encoding to generate the execution traces. We can then use the SMT encoding to further improve the DPOR technique.

References

- [1] MPI: A message-passing interface standard. <http://www.mpi-forum.org/docs/mpi-2.2/index.htm>

- [2] The multicore association. <http://www.multicore-association.org>
- [3] The multicore association resource management API. <http://www.multicore-association.org/workgroup/mcapi.php>
- [4] Bronevetsky, G.: Communication-sensitive static dataflow for parallel message passing applications. In: CGO. IEEE Computer Society (2009)
- [5] Burckhardt, S., Alur, R., Martin, M.M.K.: Checkfence: Checking consistency of concurrent data types on relaxed memory models. In: ACM SIGPLAN PLDI. San Diego, California, USA (June 10–13, 2007)
- [6] Dubrovin, J., Junttila, T., Heljanko, K.: Exploiting step semantics for efficient bounded model checking of asynchronous systems. In: Science of Computer Programming. pp. 77(10–11):1095–1121 (2012)
- [7] Dutertre, B., de Moura Leonardo: A fast linear-arithmetic solver for DPLL(T). In: CAV. vol. 4144 of LNCS, pp. 81–94. Springer-Verlag (2006)
- [8] Elwakil, M., Yang, Z.: CRI: Symbolic debugger for mcapi applications. In: Automated Technology for Verification and Analysis (2010)
- [9] Elwakil, M., Yang, Z.: Debugging support tool for mcapi applications. In: PADTAD '10: Proceedings of the 8th Workshop on Parallel and Distributed Systems (2010)
- [10] Flanagan, C., Godefroid, P.: Dynamic partial-order reduction for modeling checking software. In: POPL. pp. 110–121. ACM Press, New York, NY, USA (2005)
- [11] Gray, I., Audsley, N.: Targeting complex embedded architectures by combining the multicore communications API (mcapi) with compile-time virtualisation. In: LCTES. ACM, Chicago, Illinois, USA (2011)
- [12] Kahlon, V., Wang, C., Gupta, A.: Monotonic partial order reduction: An optimal symbolic partial order reduction technique. In: ACM CAV. pp. 398–413. Springer, Berlin/Heidelberg, Grenoble, France (June 26–July 02, 2009)
- [13] Lahiri, S.: SMT-based modular analysis of sequential systems code. In: CAV. Springer-Verlag (2011)
- [14] Lahiri, S., Qadeer, S.: Back to the future revisiting precise program verification using SMT solvers. In: POPL. ACM, San Francisco, California, USA (2008)
- [15] Morse, E., Vrvilo, N., Mercer, E., McCarthy, J.: Modeling asynchronous message passing for C program. In: VMCAI. vol. 7148 of LNCS (2012)
- [16] de Moura, L., Bjørner, N.: Z3: An efficient SMT solver. In: TACAS. vol. 4963, pp. 337–340. Springer, Heidelberg (2008)
- [17] Netzer, R., Brennan, T., Damodaran-Kamal, S.: Debugging race conditions in message-passing programs. In: ACM SIGMETRICS Symposium on Parallel and Distributed Tools. pp. 31–40. Philadelphia, PA, USA (1996)
- [18] Sharma, S., Gopalakrishnan, G., Mercer, E., Holt, J.: Mcc - a runtime verification tool for mcapi user applications. In: FMCAD (2009)
- [19] Vakkalanka, S., Vo, A., Gopalakrishnan, G., Kirby, R.: Reduced execution semantics of mpi: From theory to practice. In: FM. pp. 724–740 (2009)
- [20] Wang, C., Chaudhuri, S., Gupta, A., Yang, Y.: Symbolic pruning of concurrent program executions. In: ESEC/FSE. pp. 23–32. ACM, New York, NY, USA (2009)
- [21] Zhang, Y., Evelyn, D.: Barrier matching for programs with textually unaligned barriers. In: PPOPP. pp. 194–204. ACM, San Jose, California, USA (2007)